CMPINF 2110

Spring 2021

Homework 06 Solutions

Create a new local DBMS in Neo4j desktop. I chose to name mine hw06.

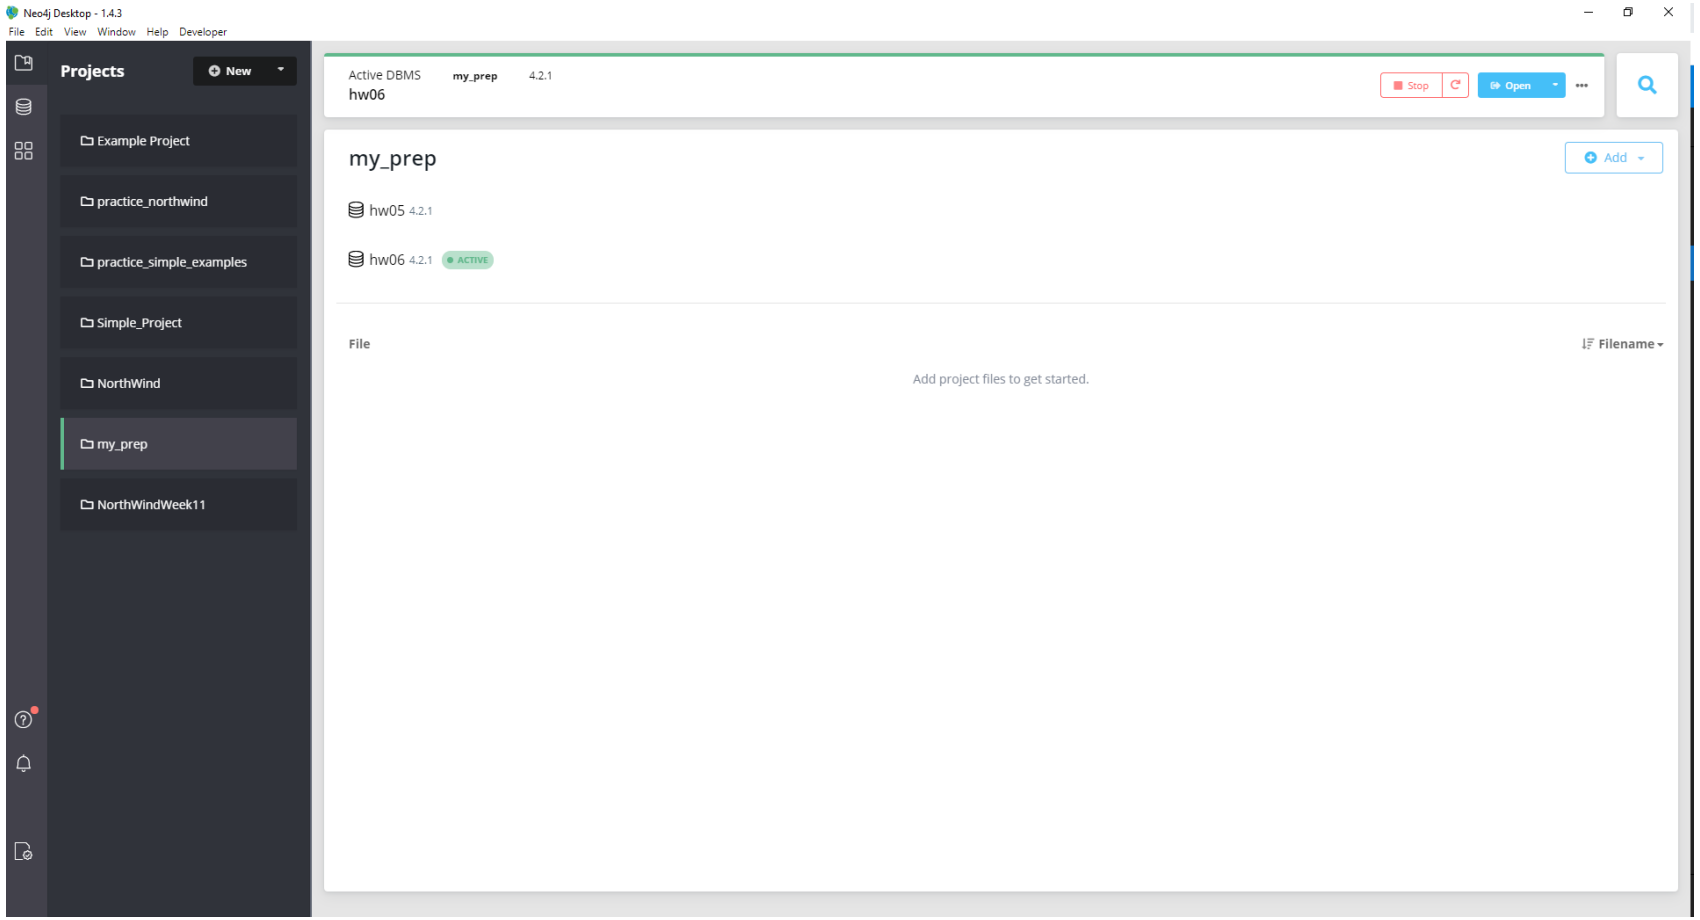

Add Device Label and create nodes with properties set to the columns of the devices table

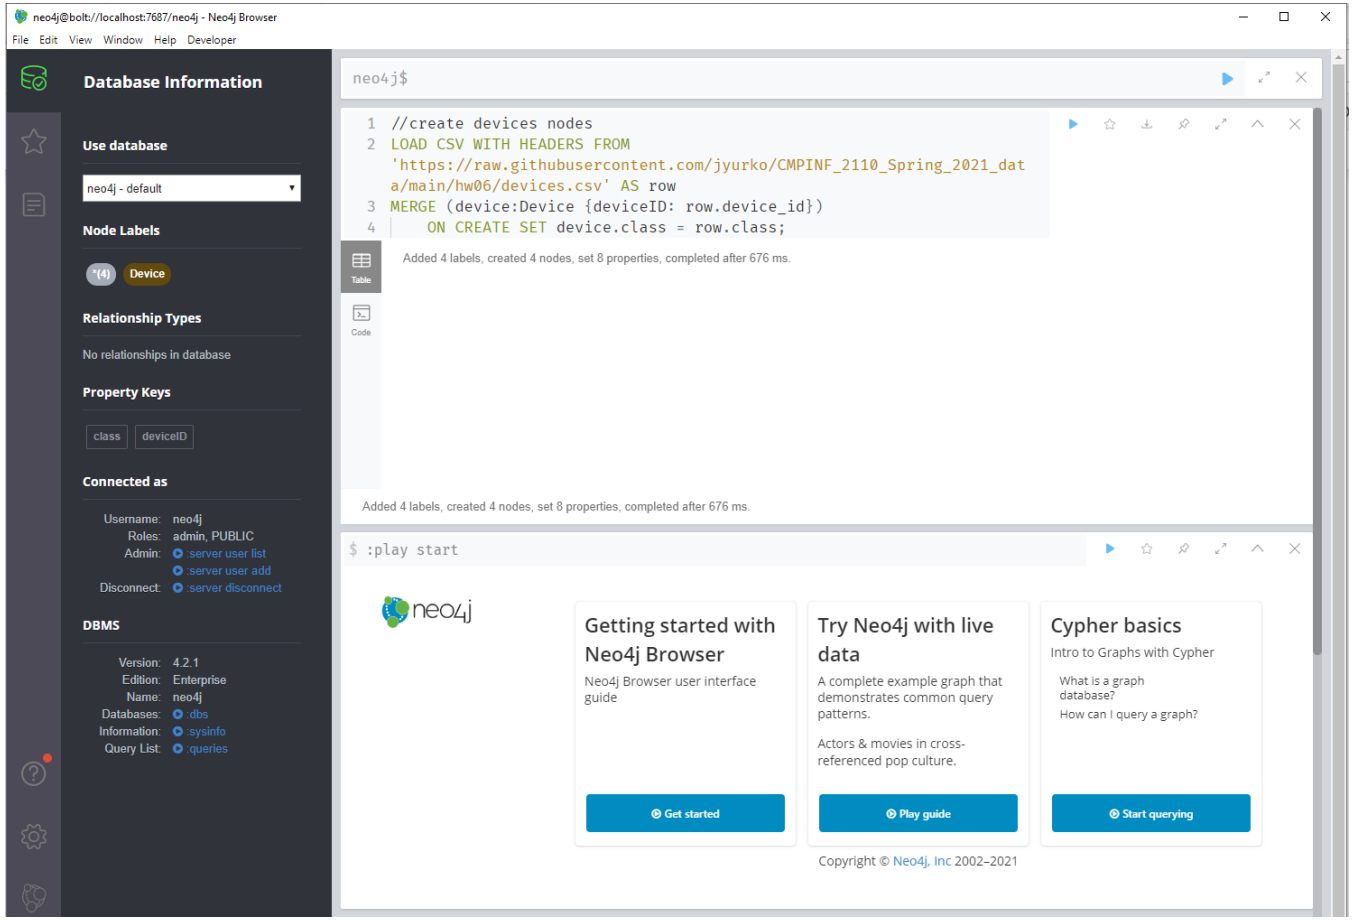

Add Employee Label and create nodes with properties set to the columns of the employees table

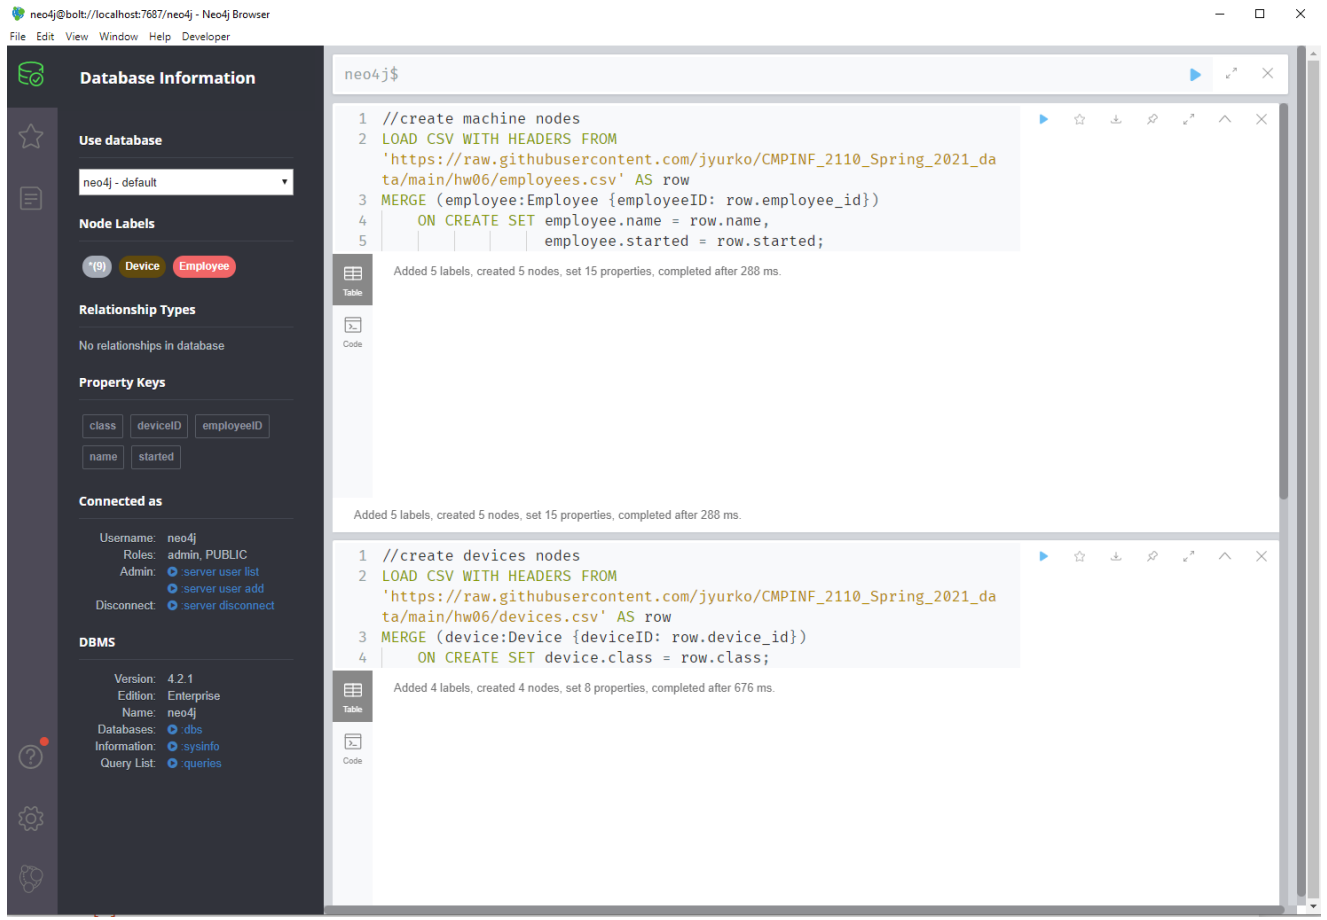

Add Machine Label and create nodes with properties set to the columns of the machines table

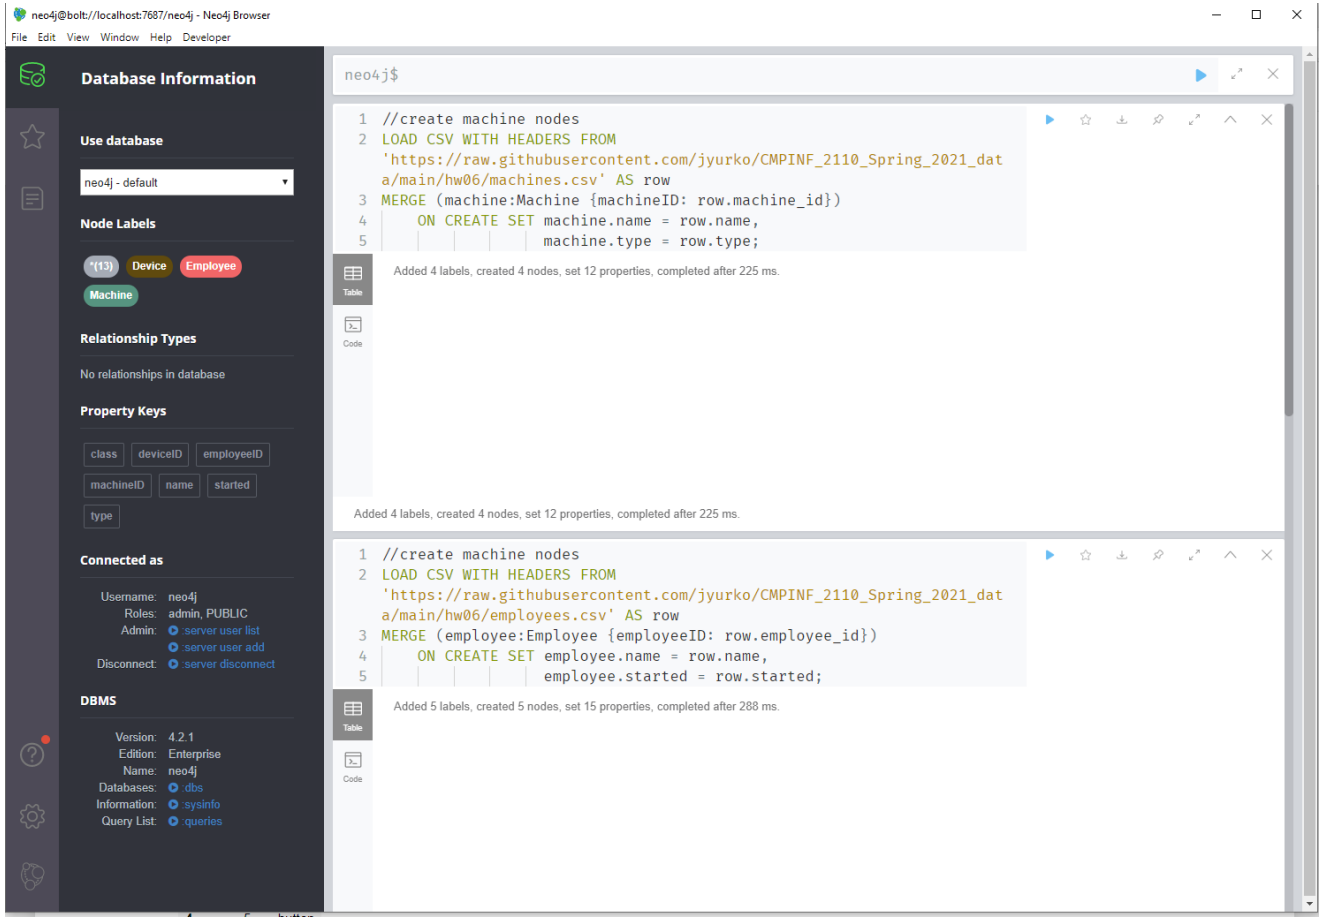

Add Part Label and create nodes with properties set to the columns of the parts table

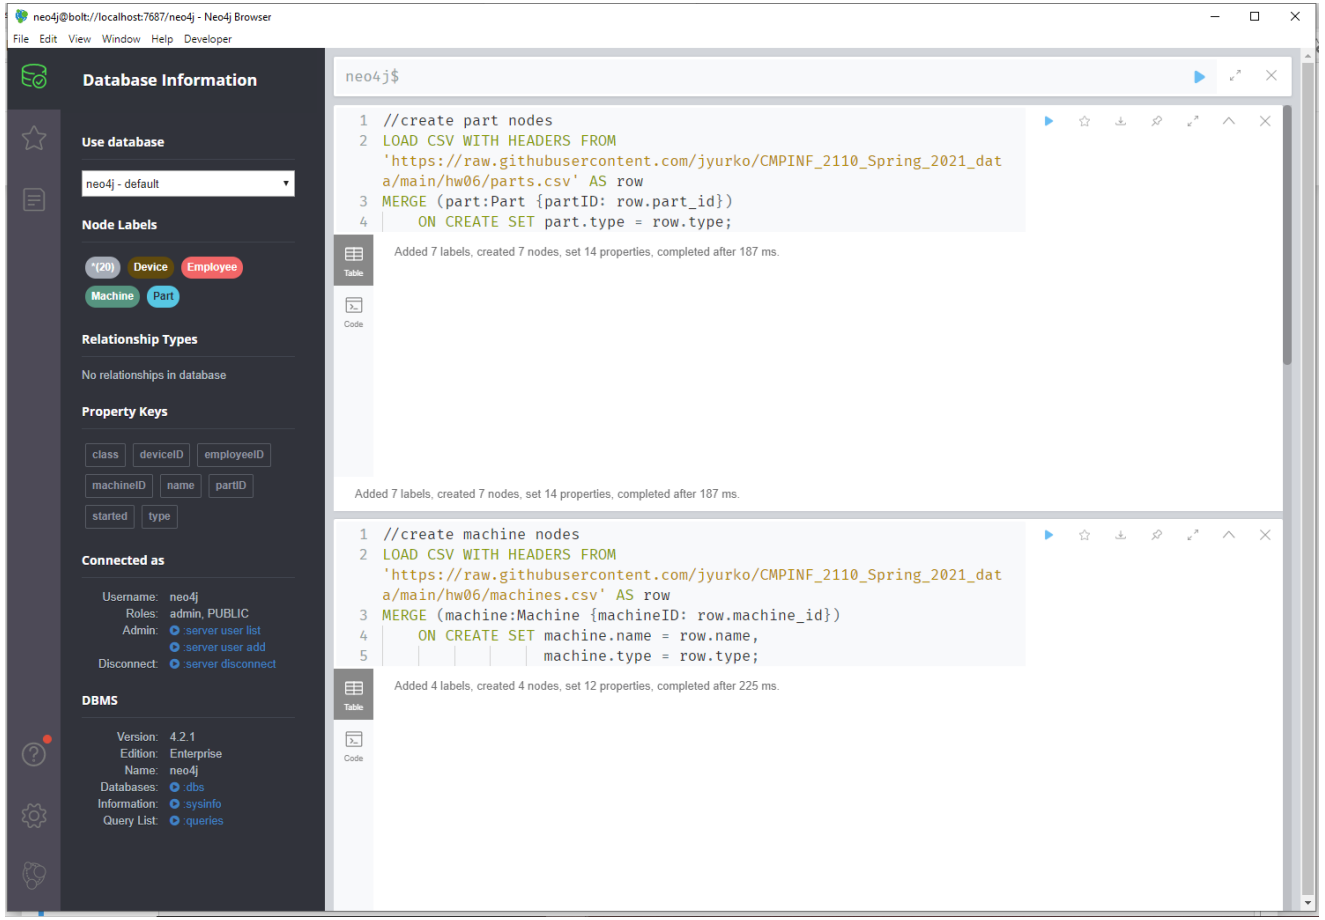

As a check, query all nodes in the graph model

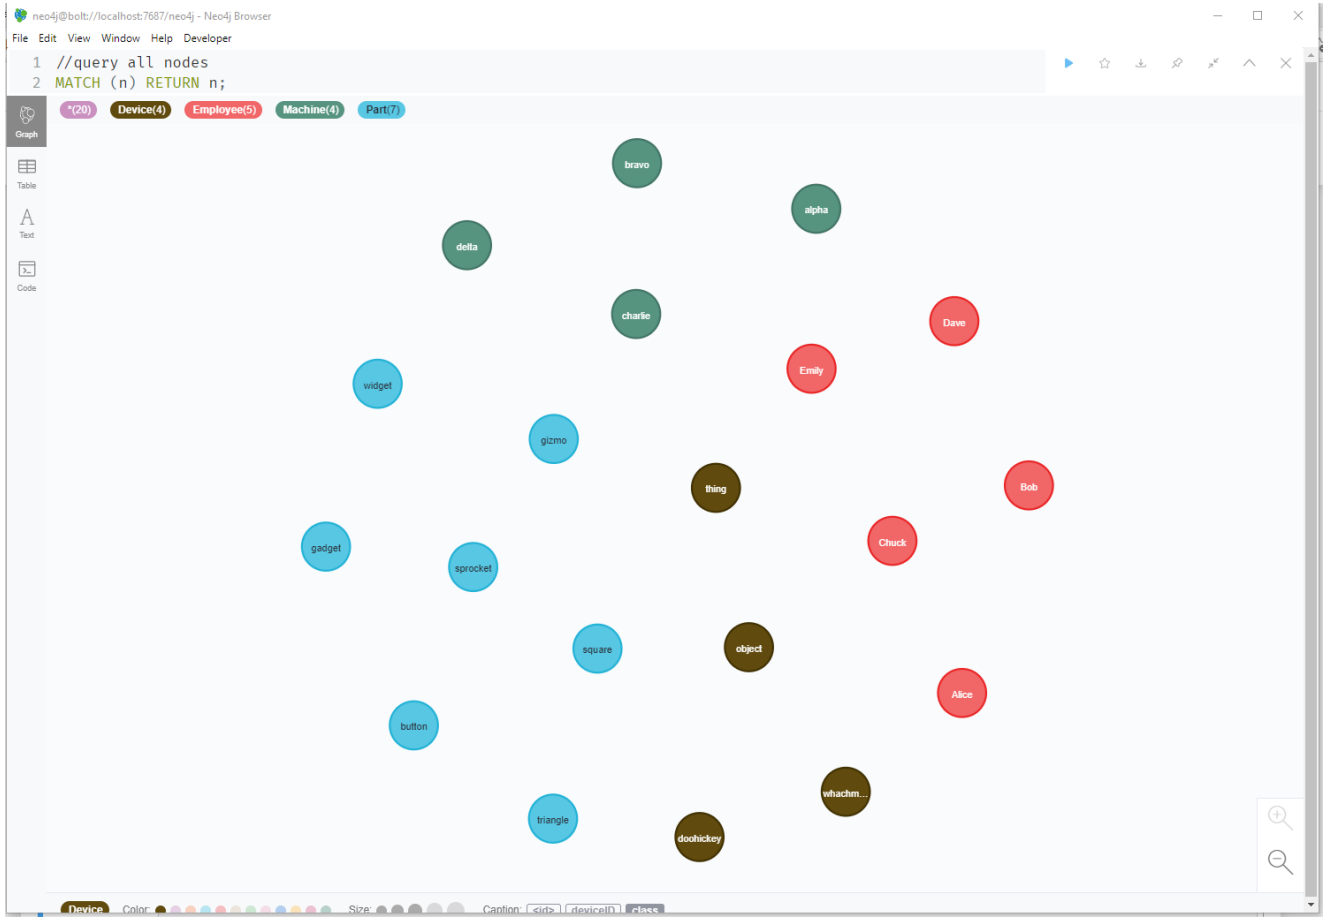

No relationships are defined as of yet!

The “meta graph” is another way to reveal no relationships are defined as of yet

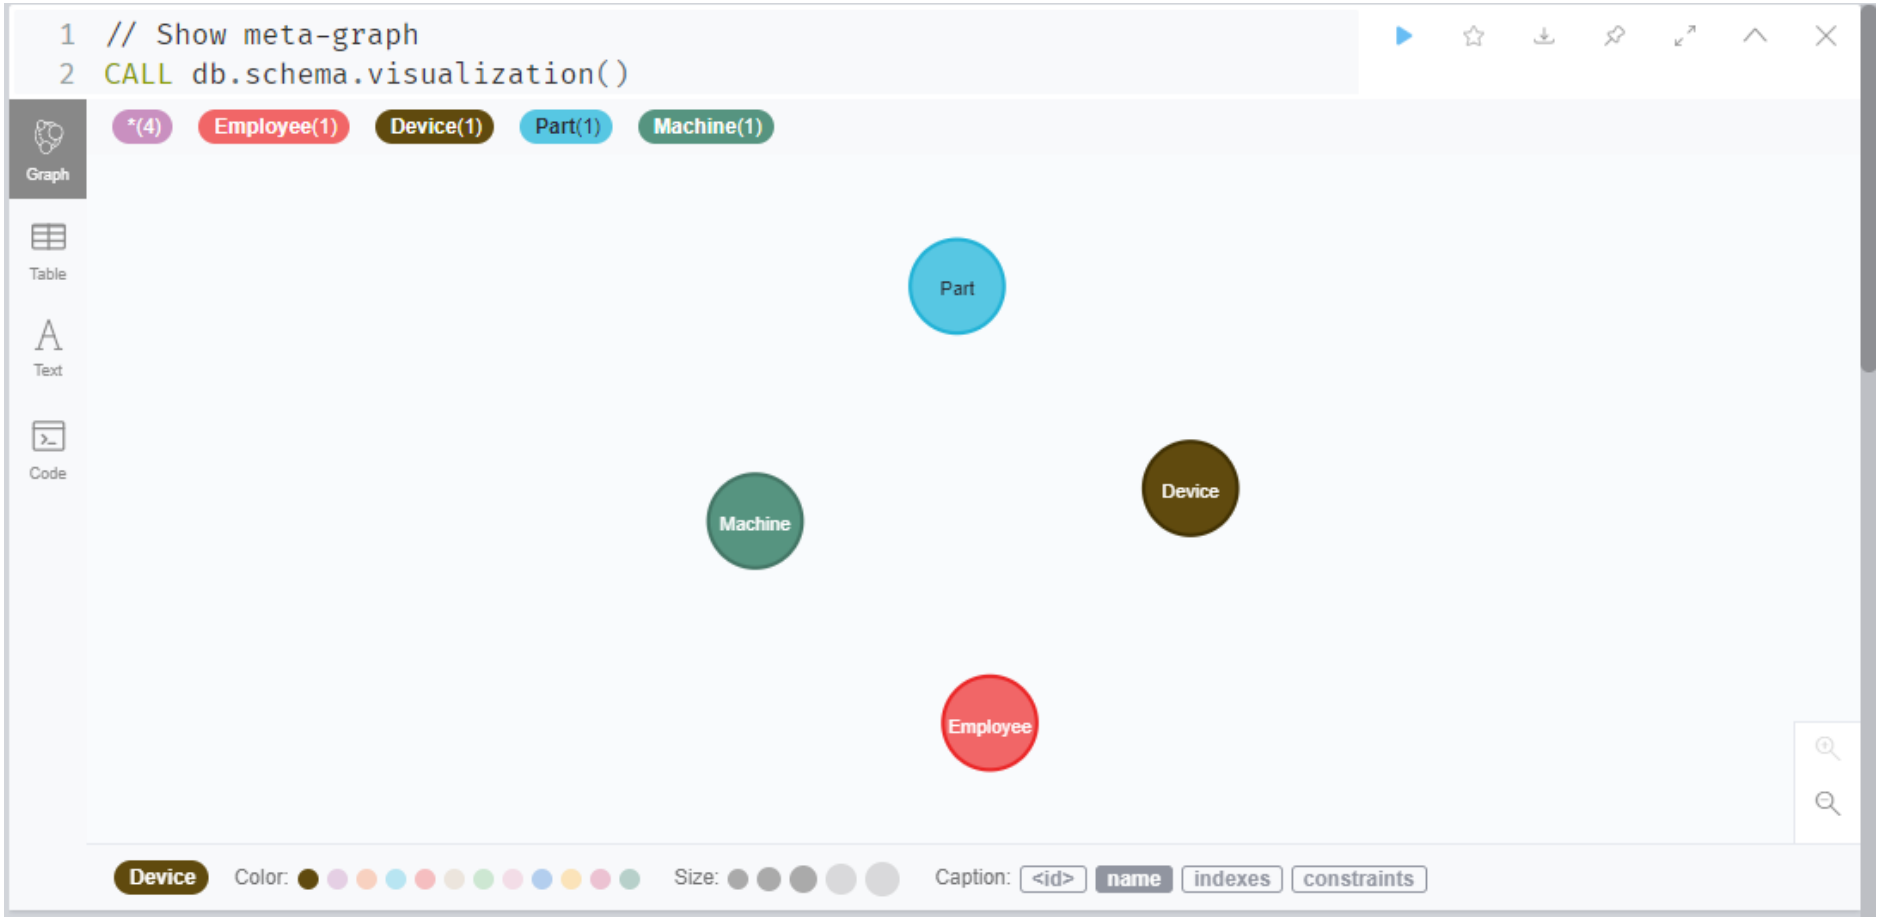

Add Job label and create nodes with jobID property

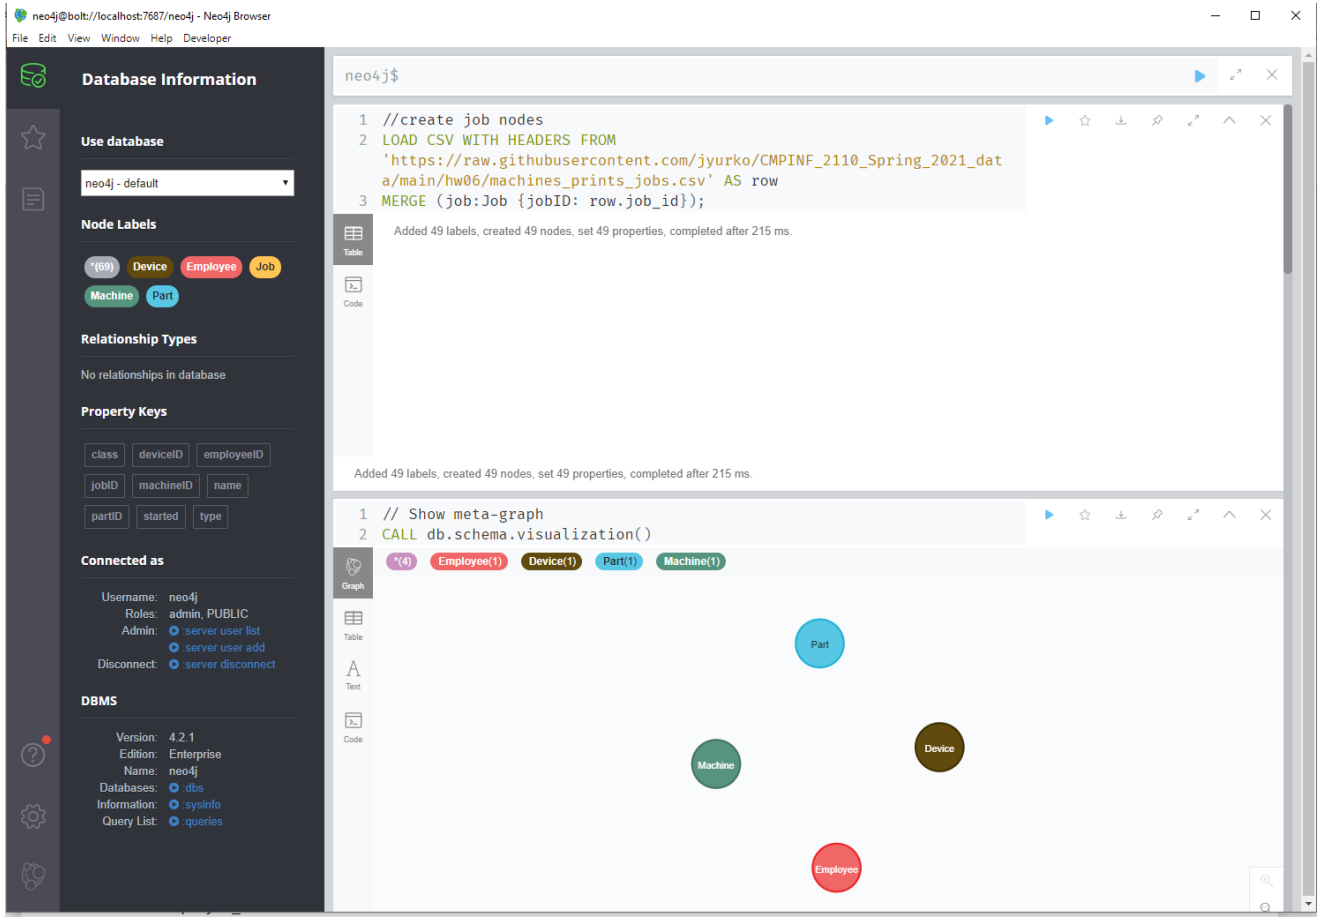

Add Item label and create nodes with itemID property

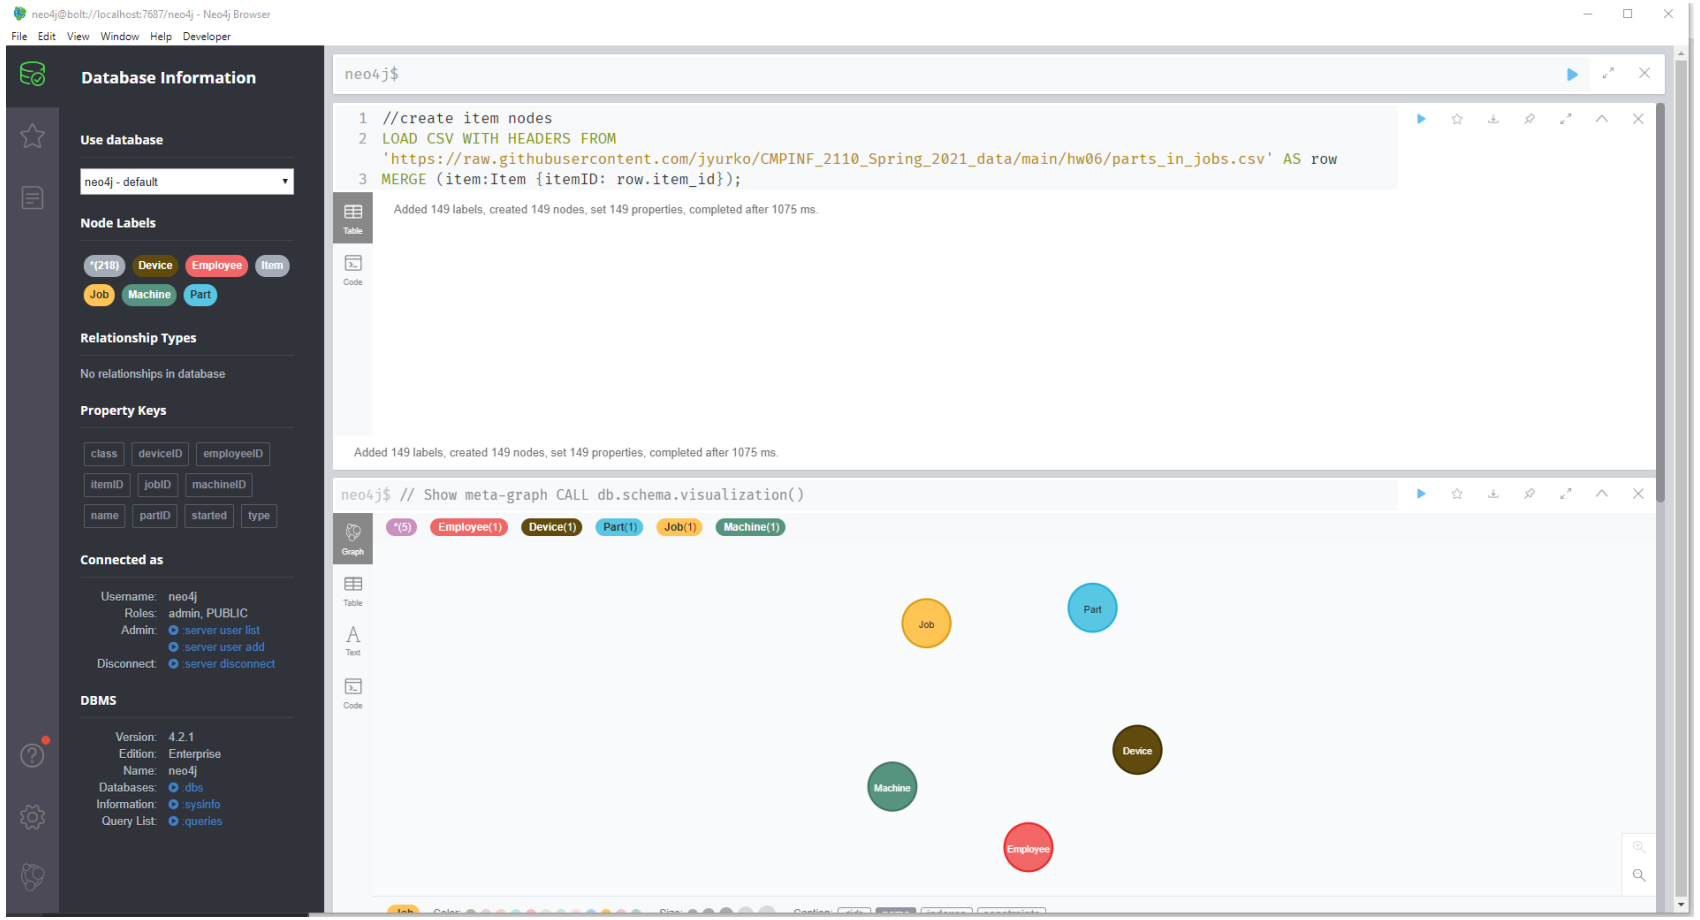

Add Assembly label and create nodes with assemblyID property

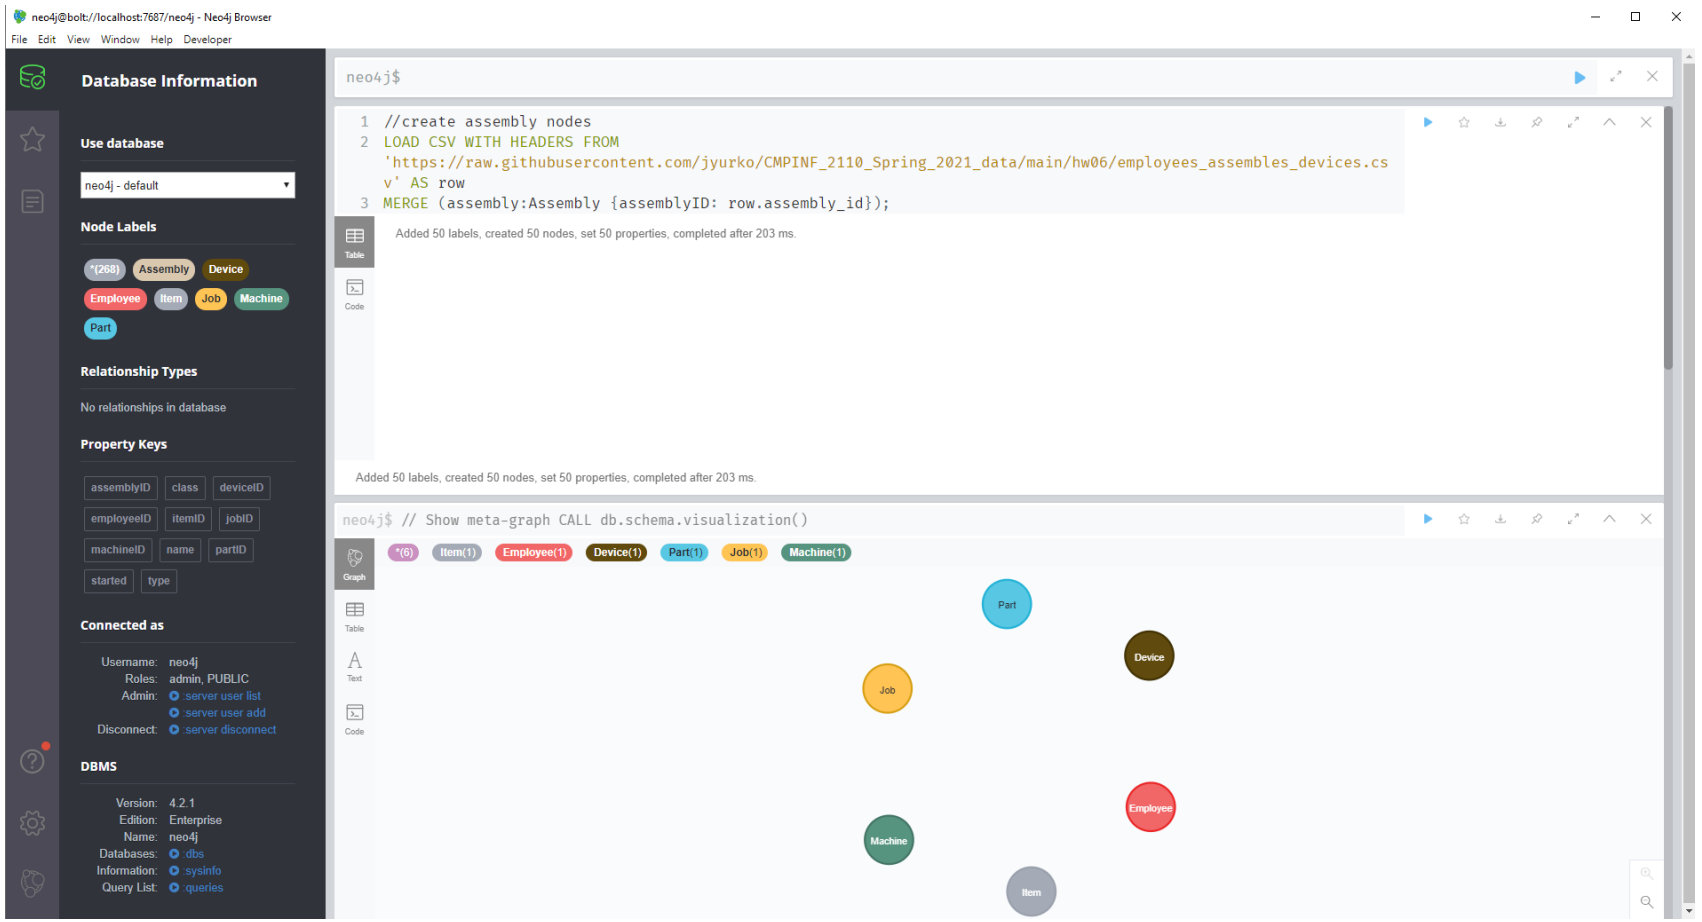

All nodes have been created. The meta-graph shows no relationships have been defined as of yet.

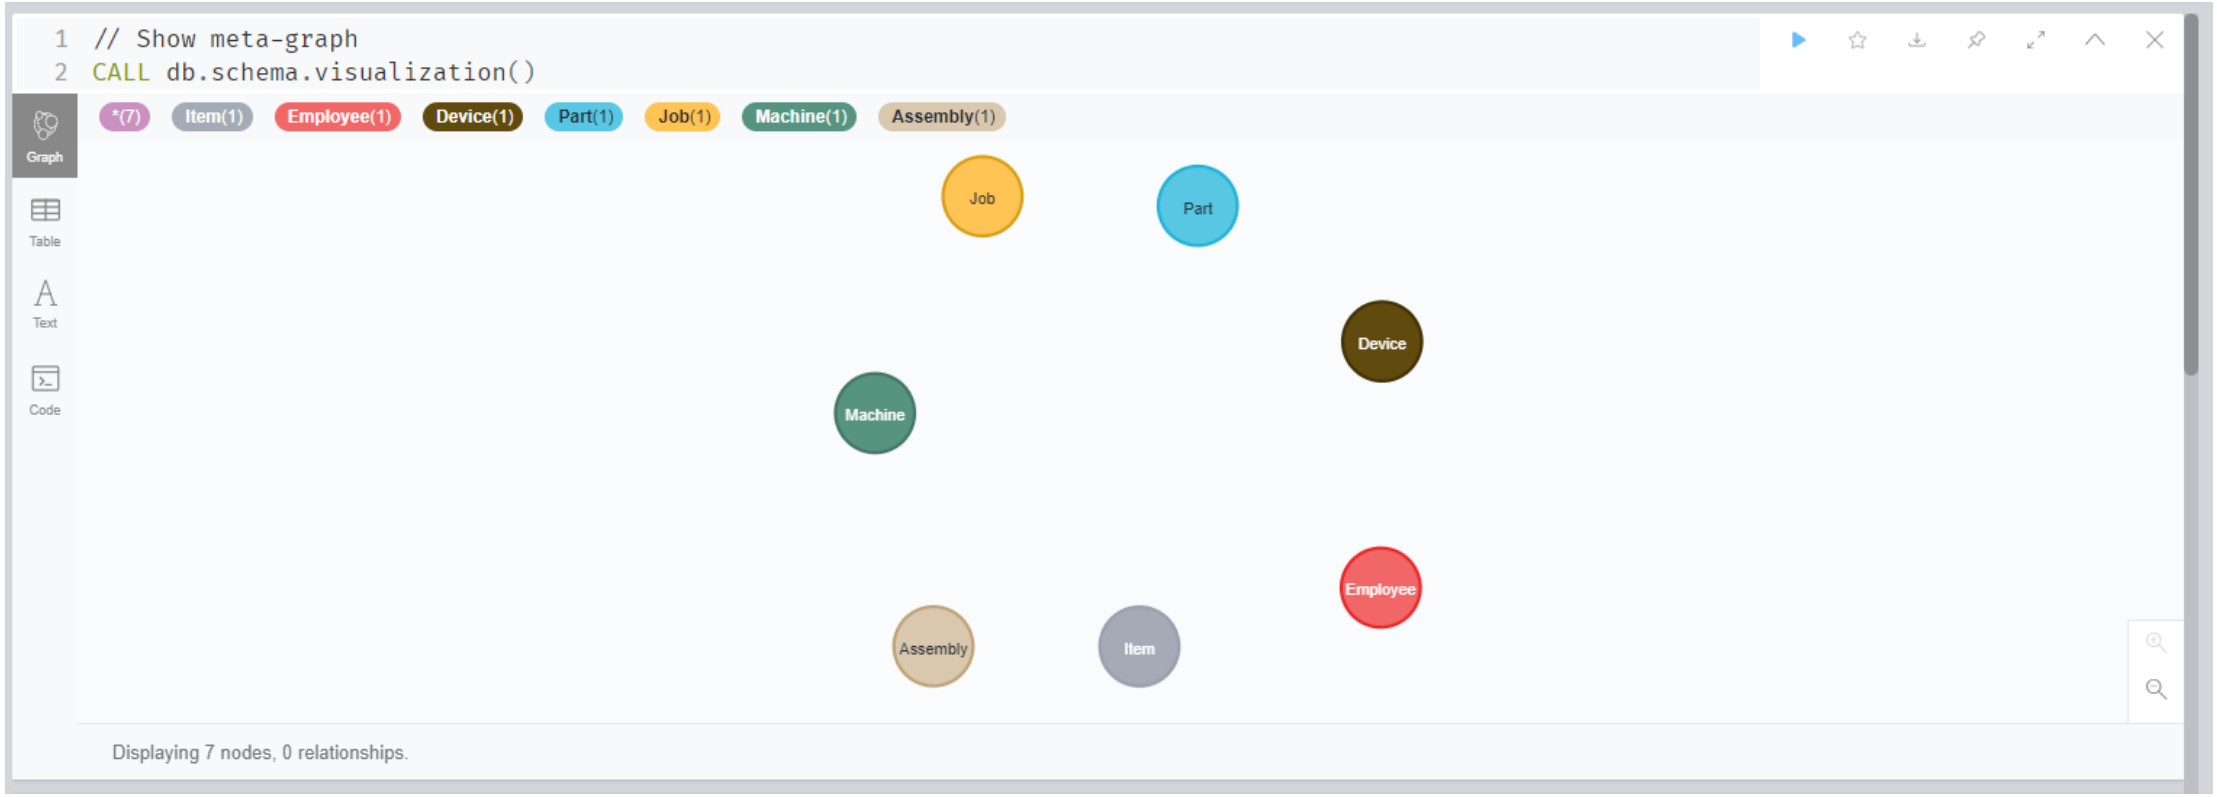

Create the ASSEMBLES relationship type between Employee nodes and Assembly nodes

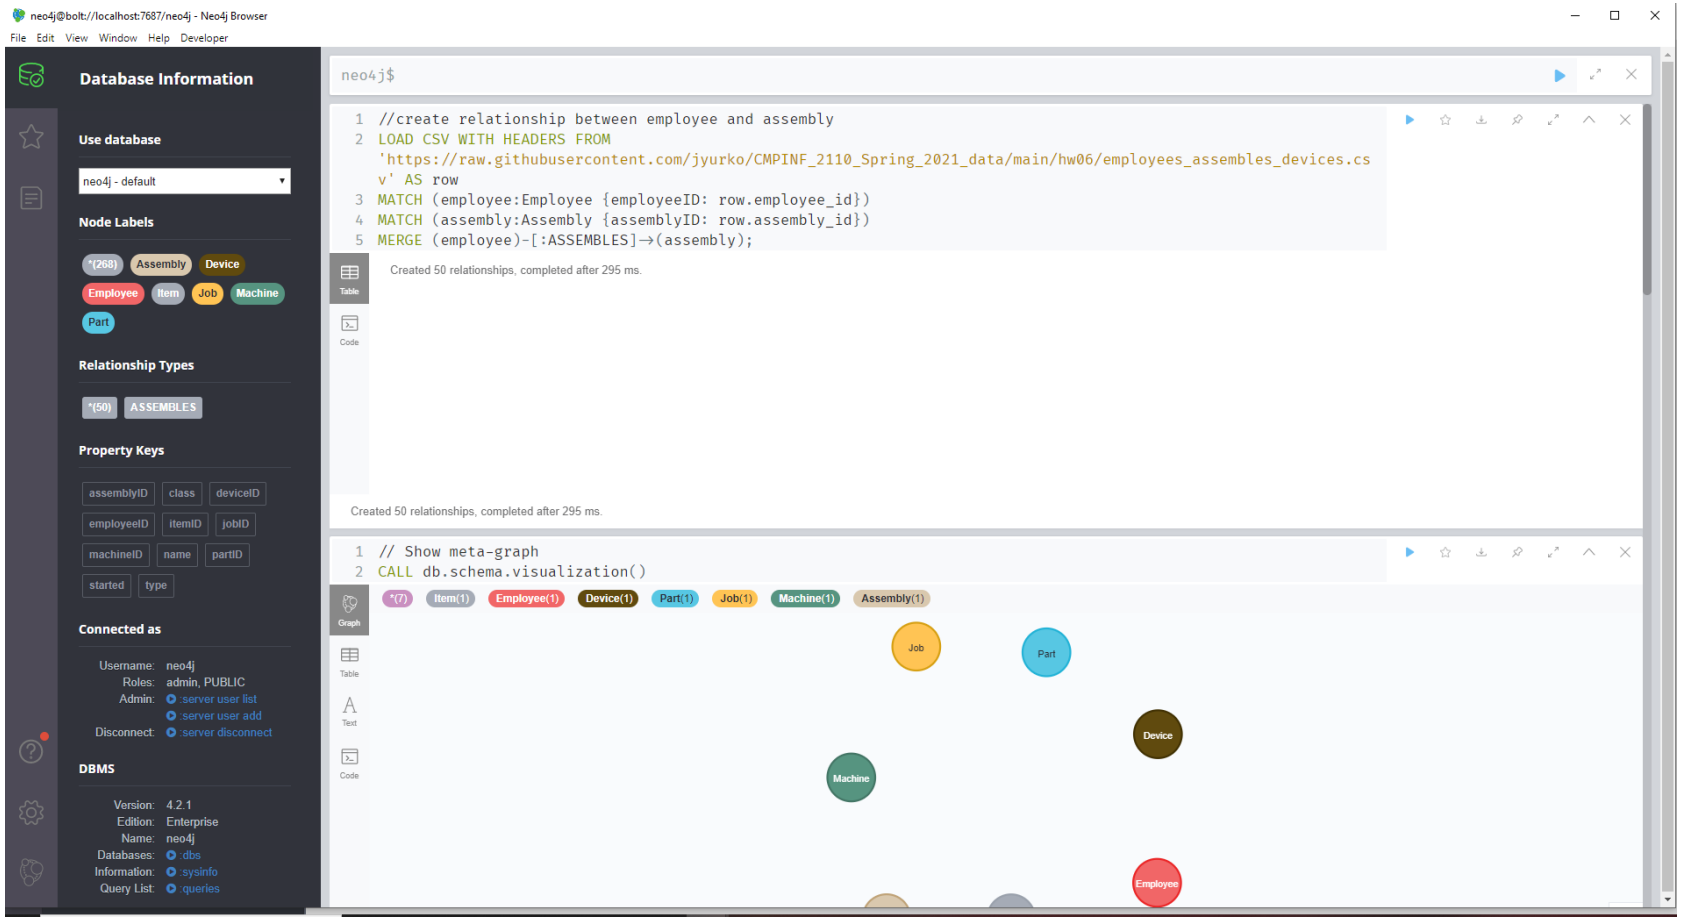

We now have a relationship!

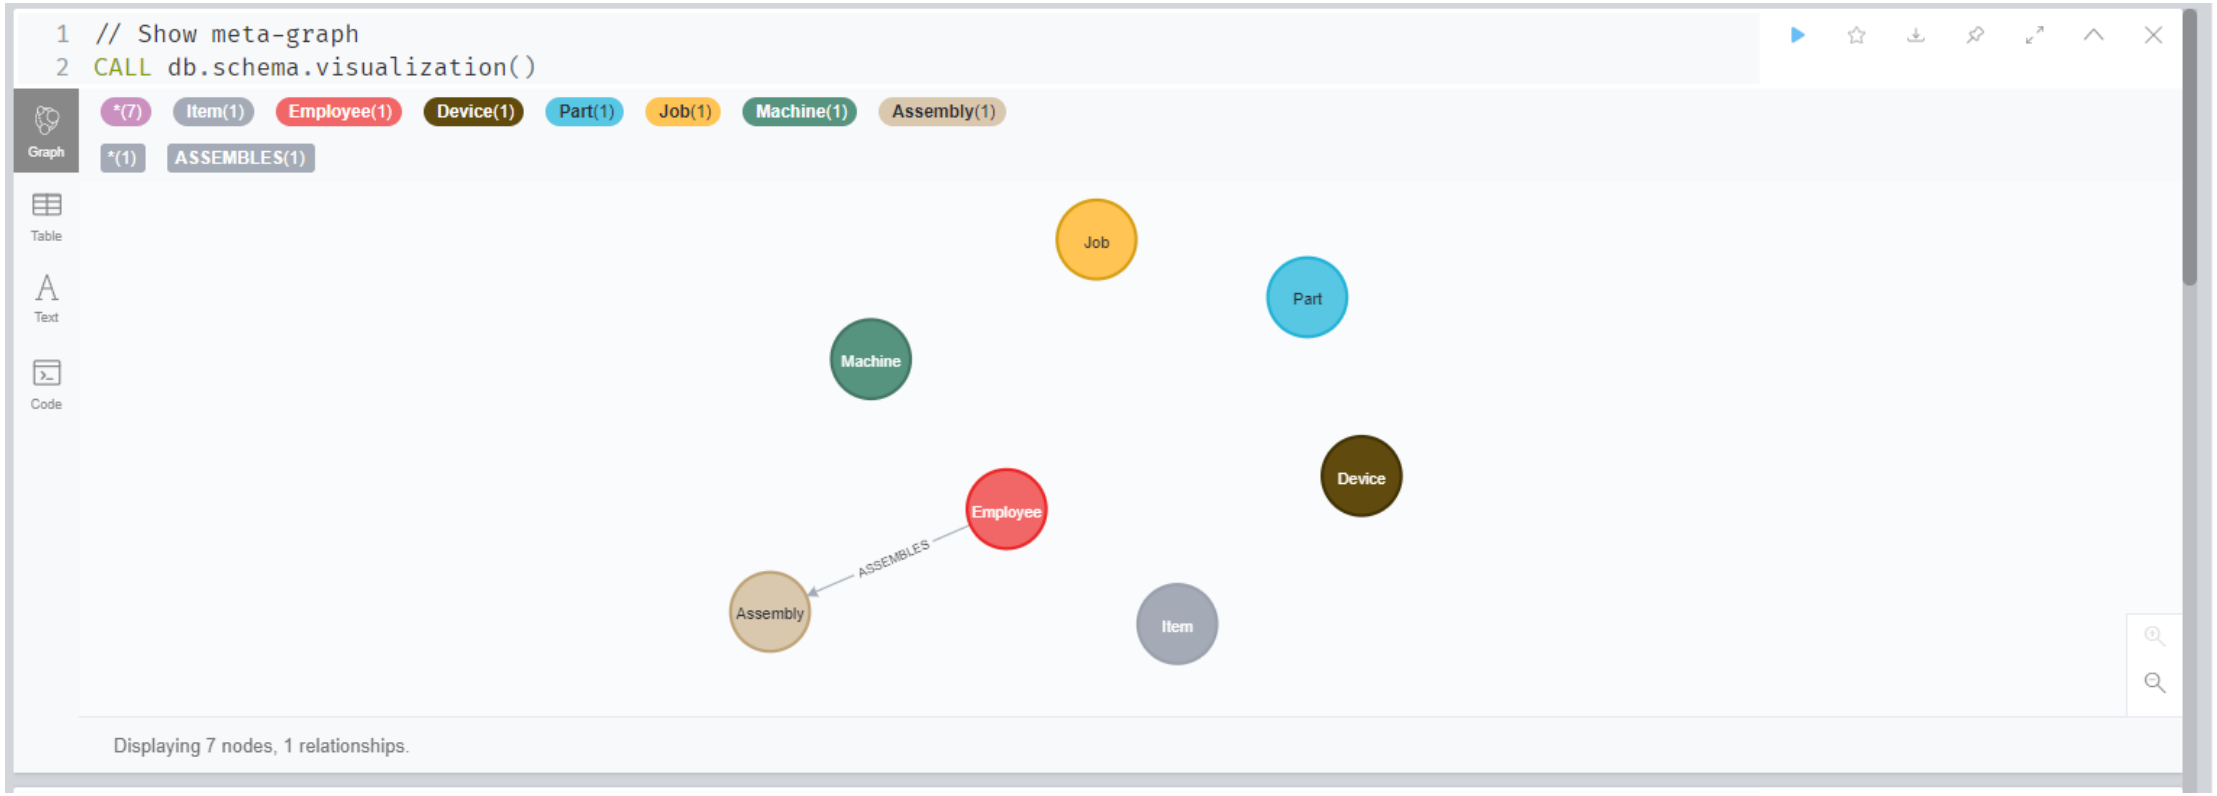

Create the OPERATES relationship type between Employee and Job nodes

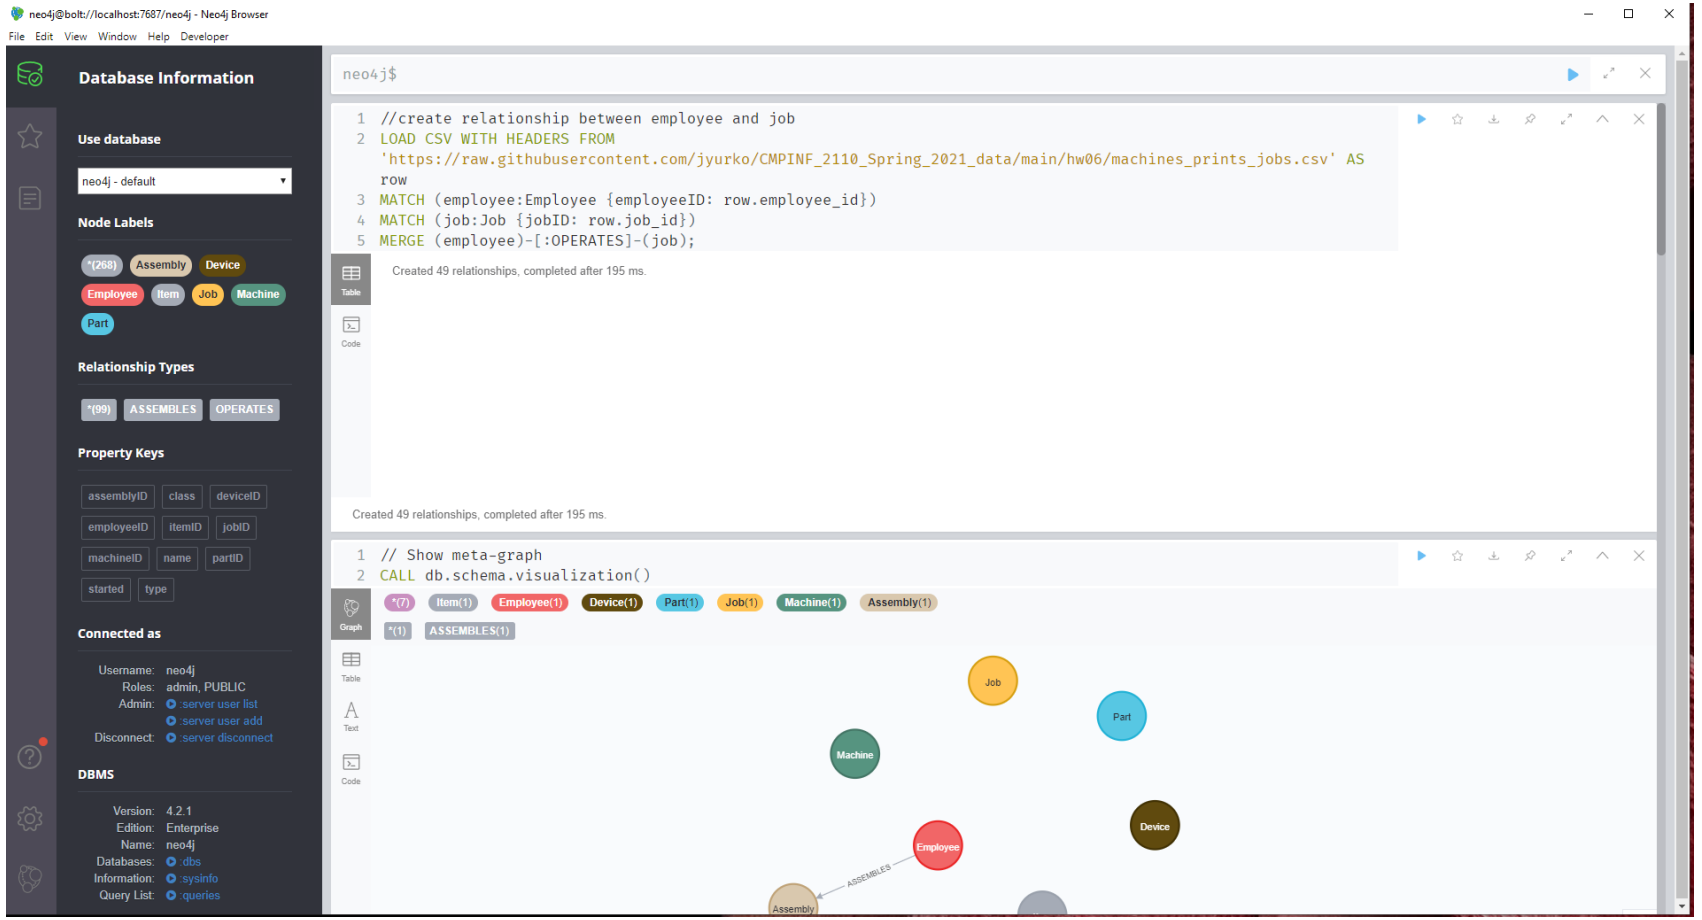

Check the meta-graph to see the two relationships

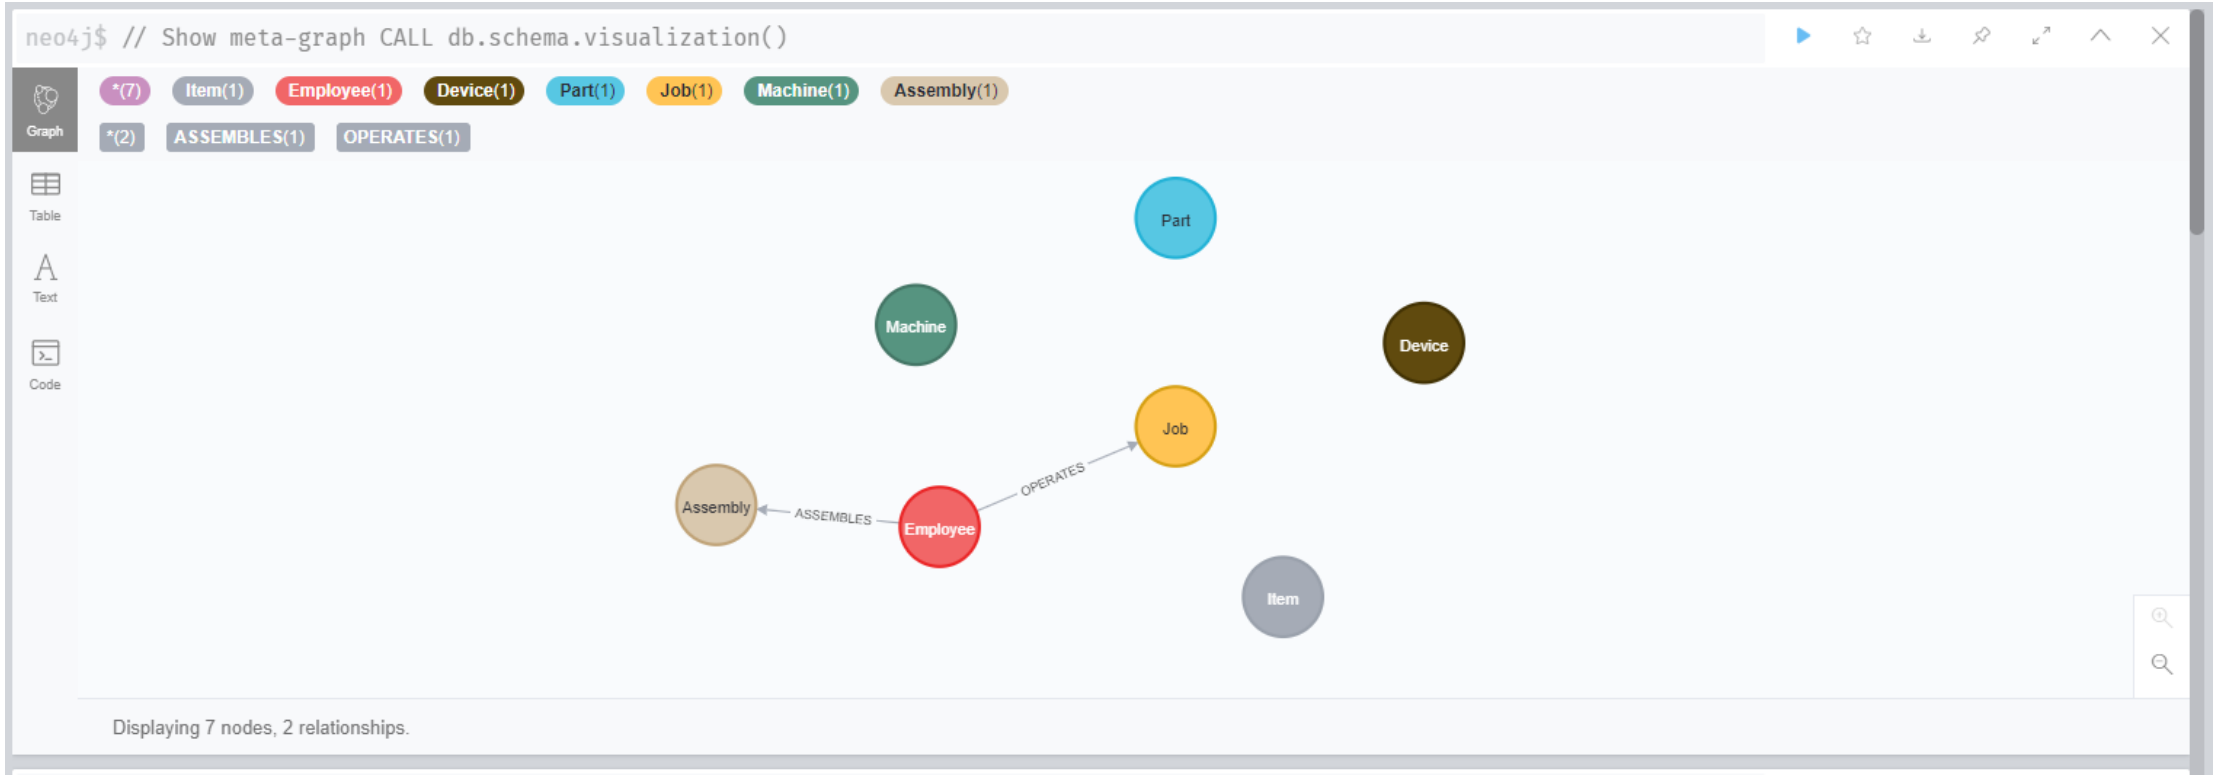

Create the PRINTS relationship type between Machine and Job nodes

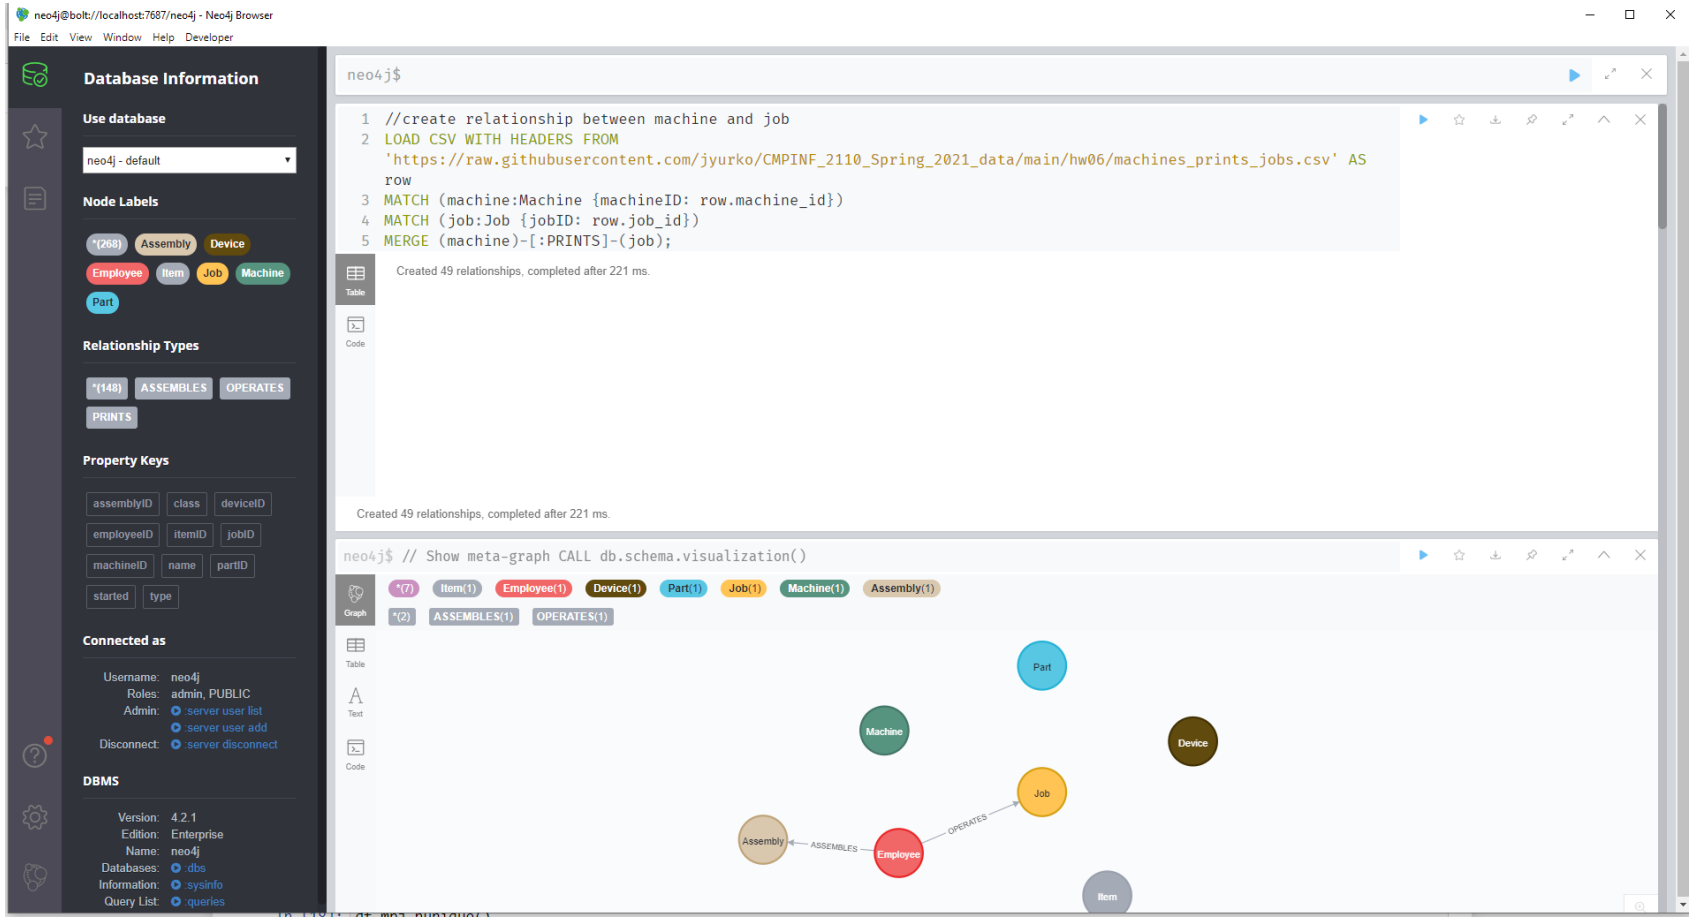

The meta-graph shows the 3 relationships currently in the graph database

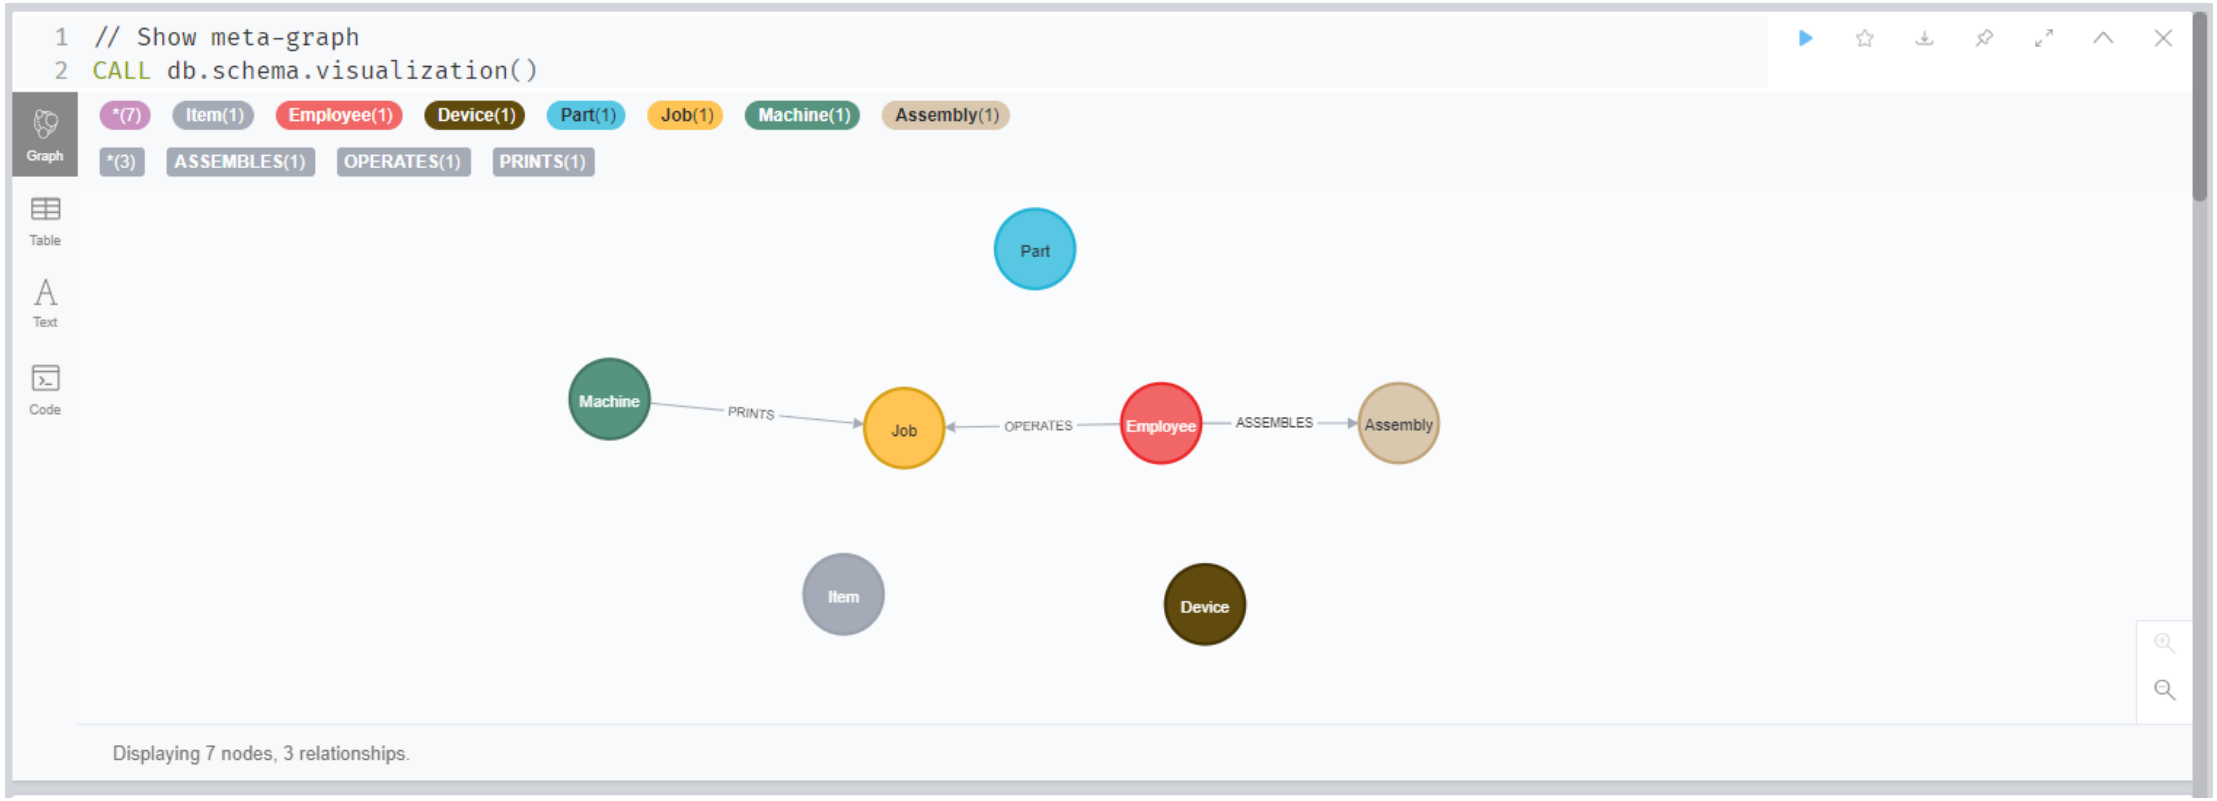

Create the IS_IN relationship type between Item and Job nodes

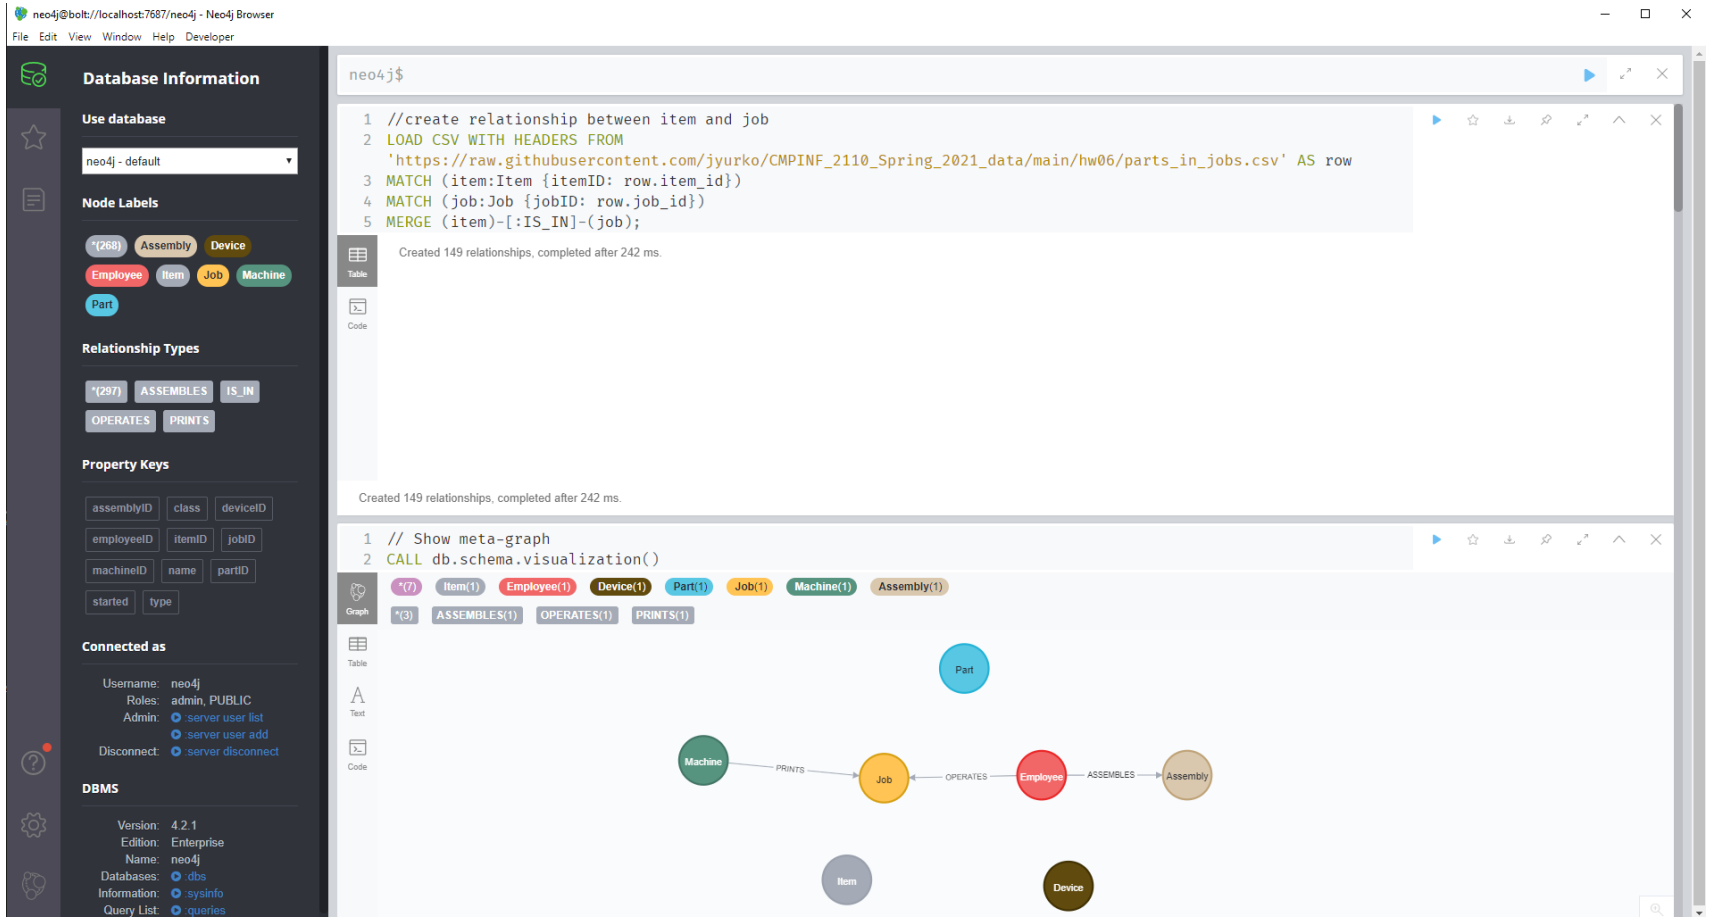

Meta-graph shows all the relationships associated with the Job Label

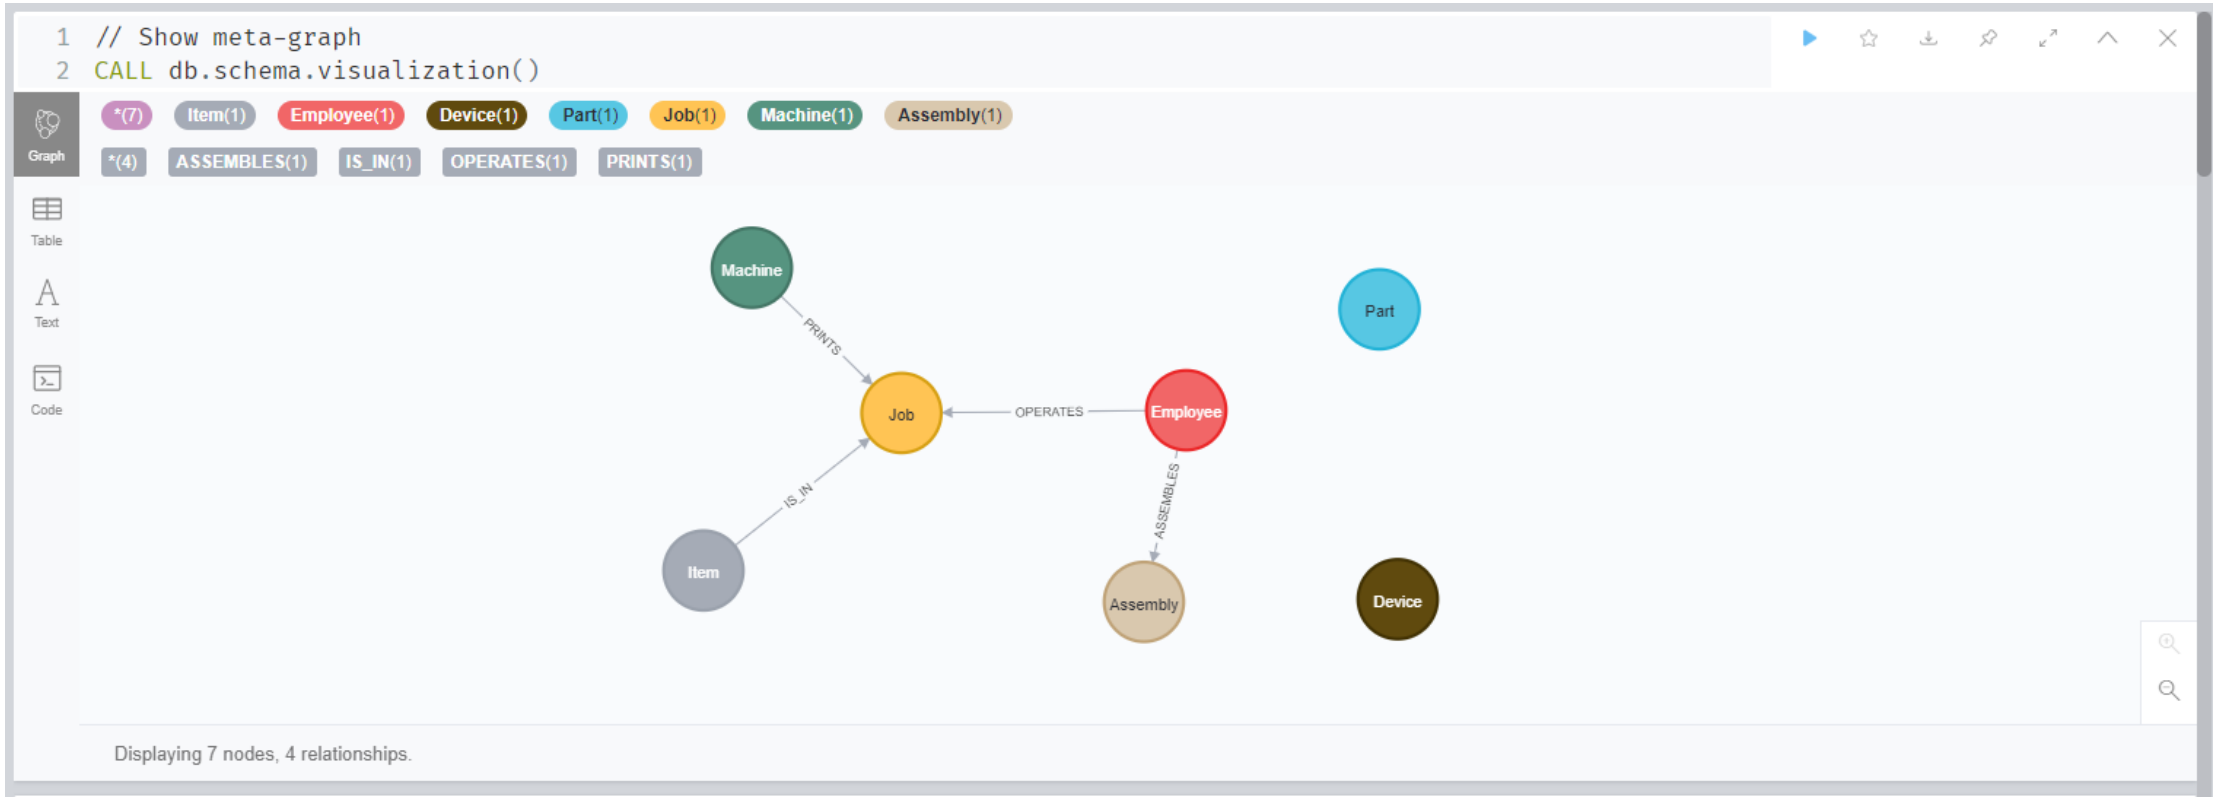

Create the IS_COMPONENT_OF relationship type between Item and Assembly nodes

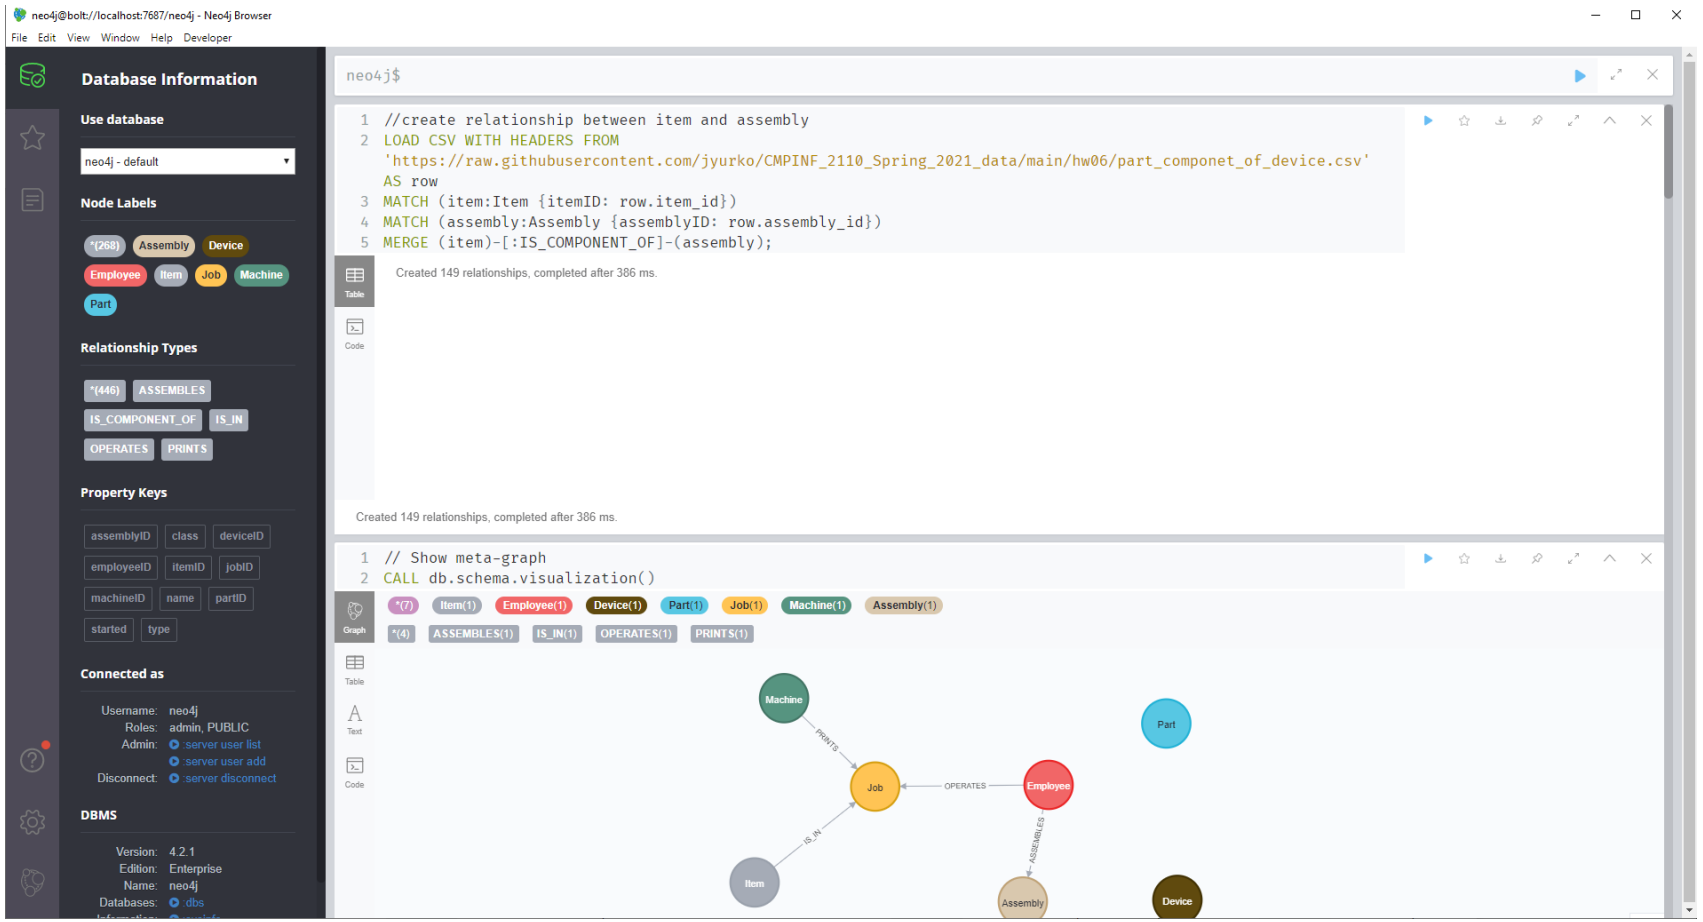

The meta-graph shows just the Part and Device labels are not connected in the network

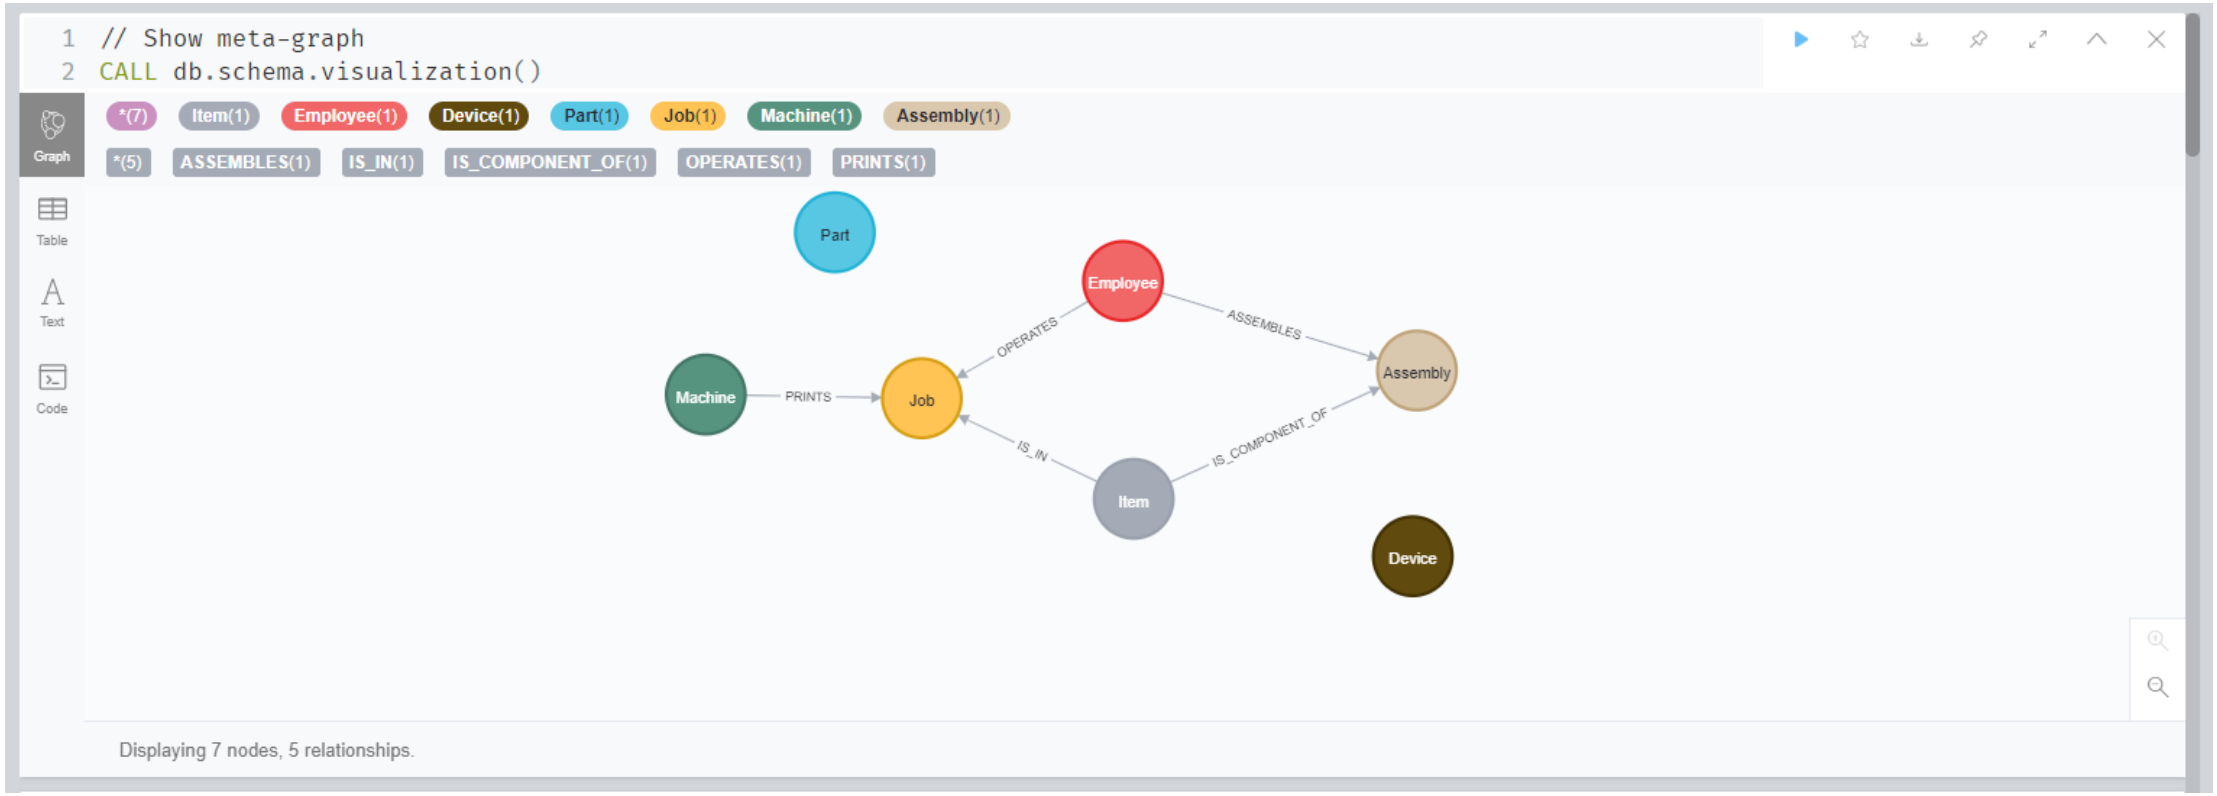

Create the IS_TYPE_OF relationship between the Item and Part nodes

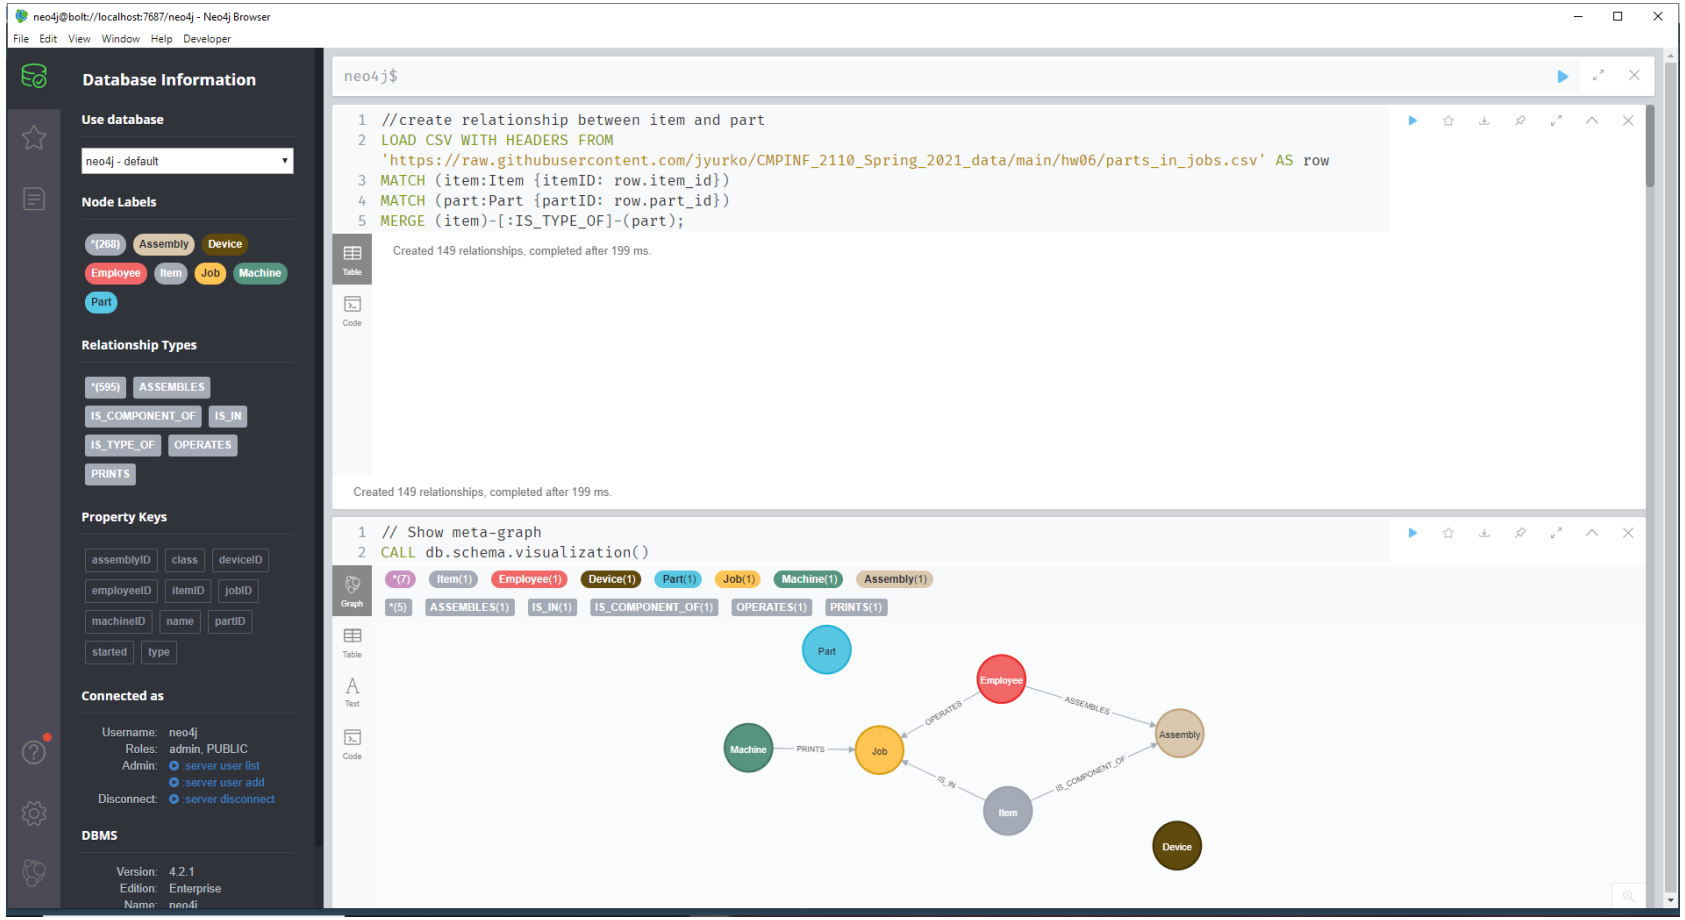

Meta-graph shows only the Device label is not connected to other Labels

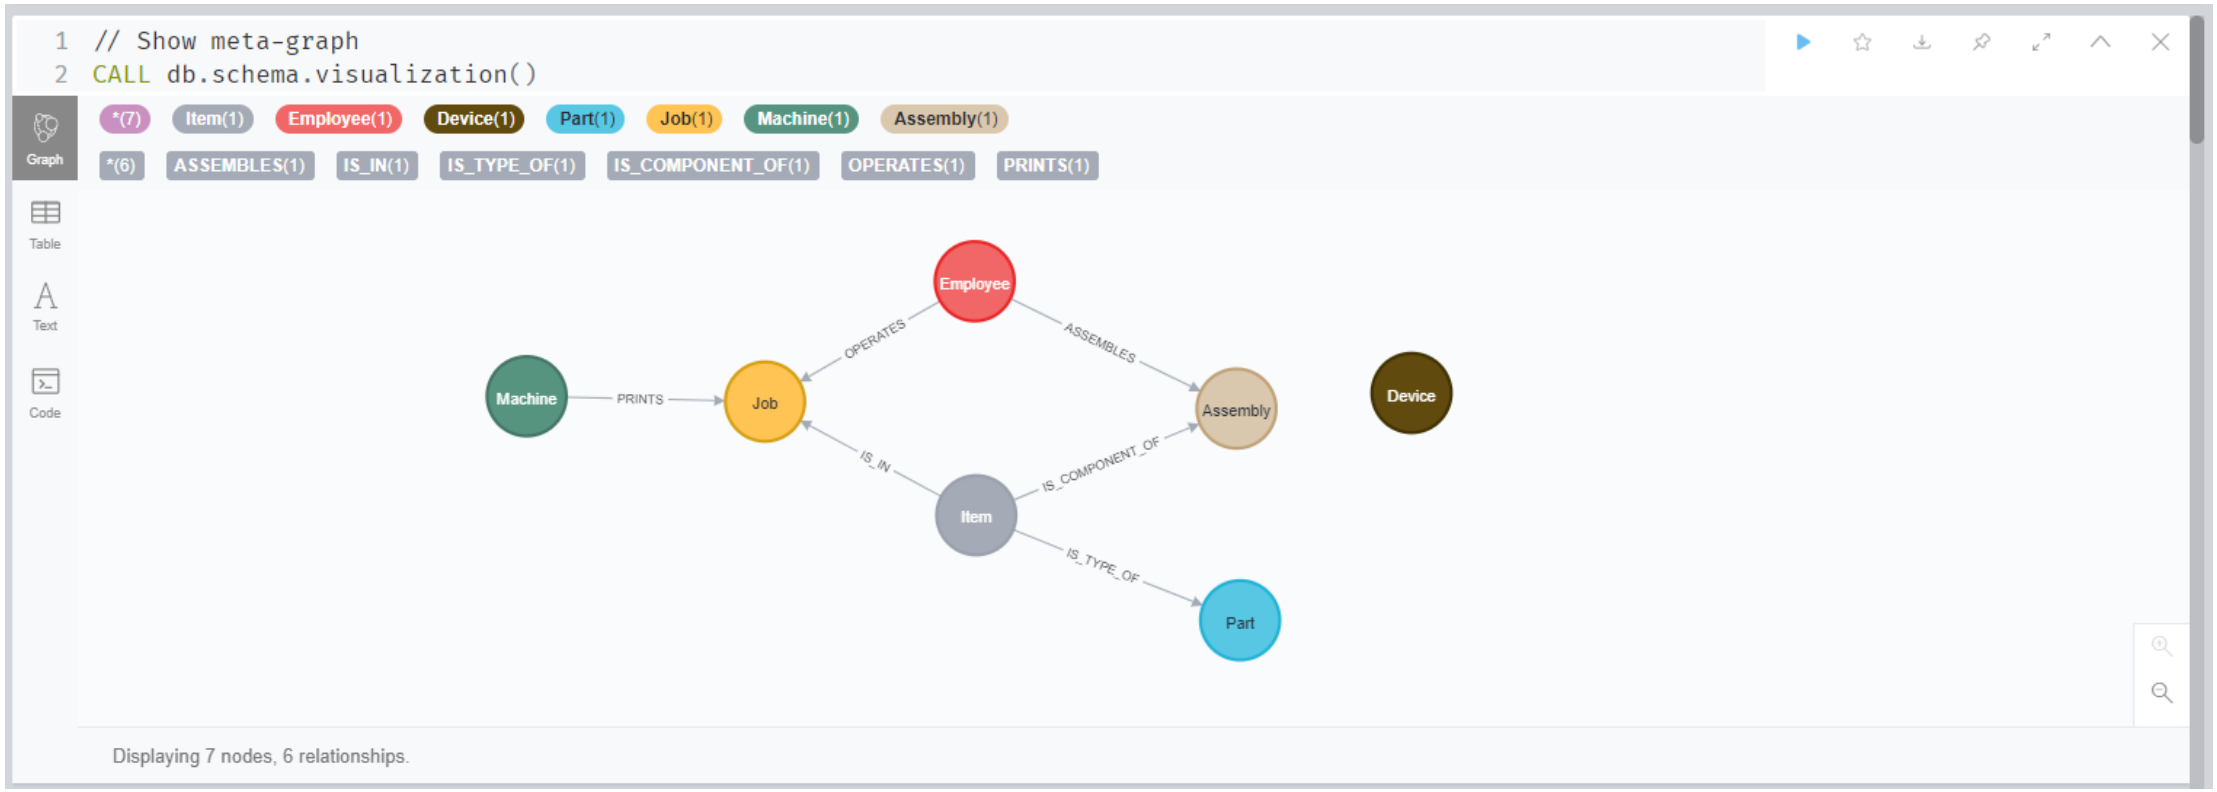

Create the IS_CLASS_OF relationship between the Assembly and Device nodes

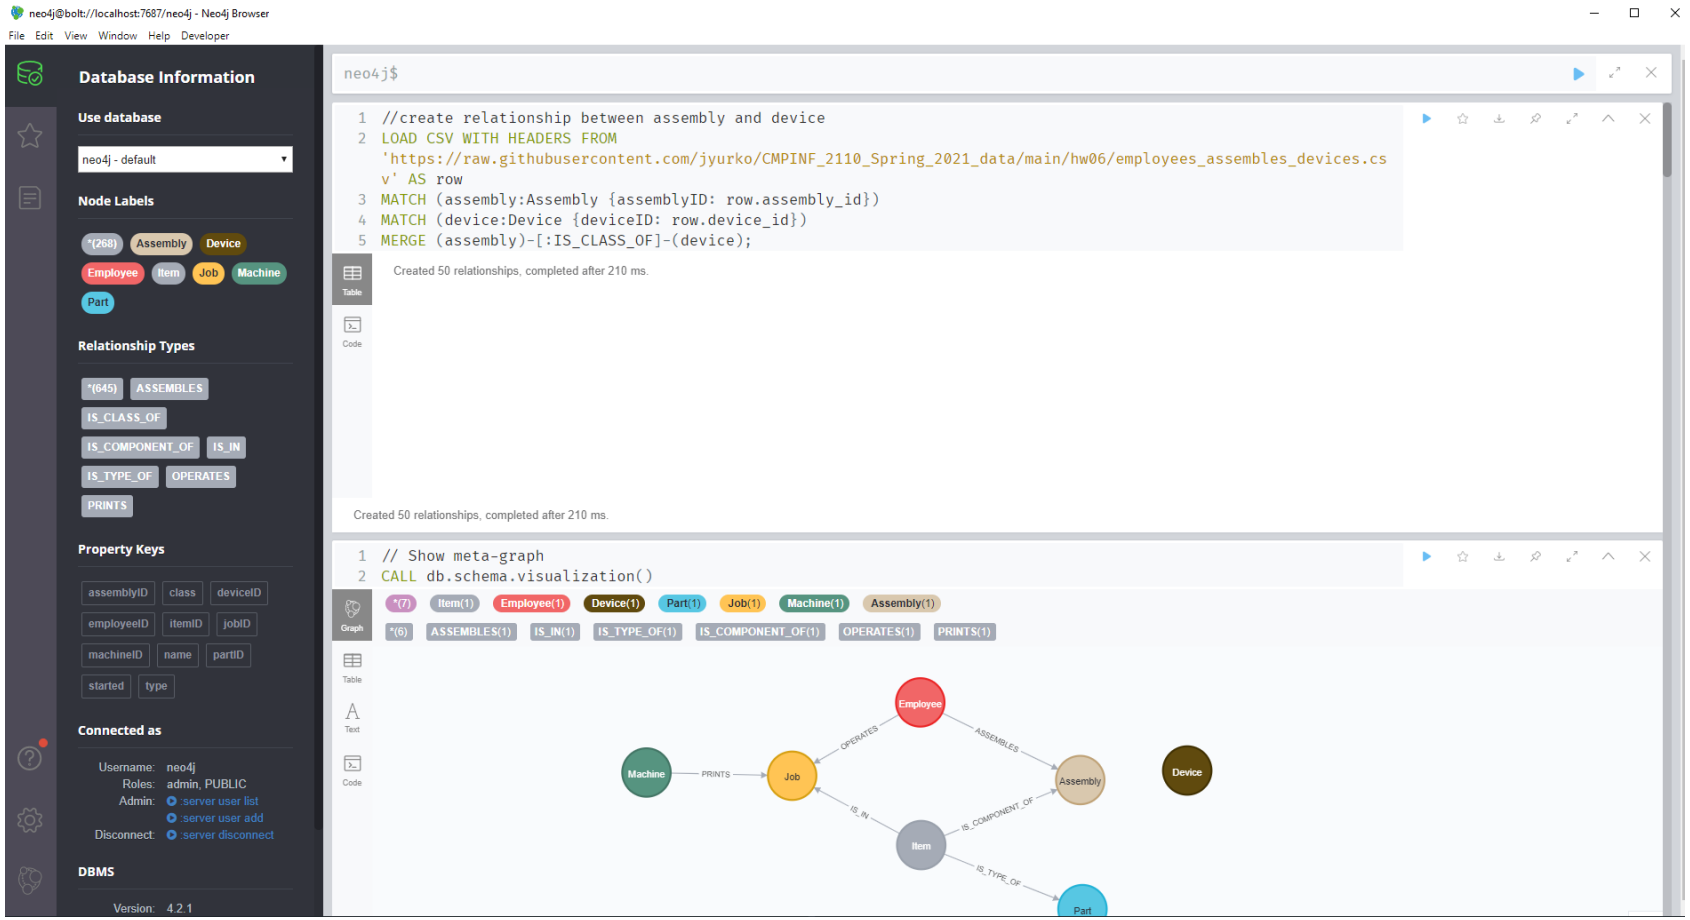

Completed graph schema!

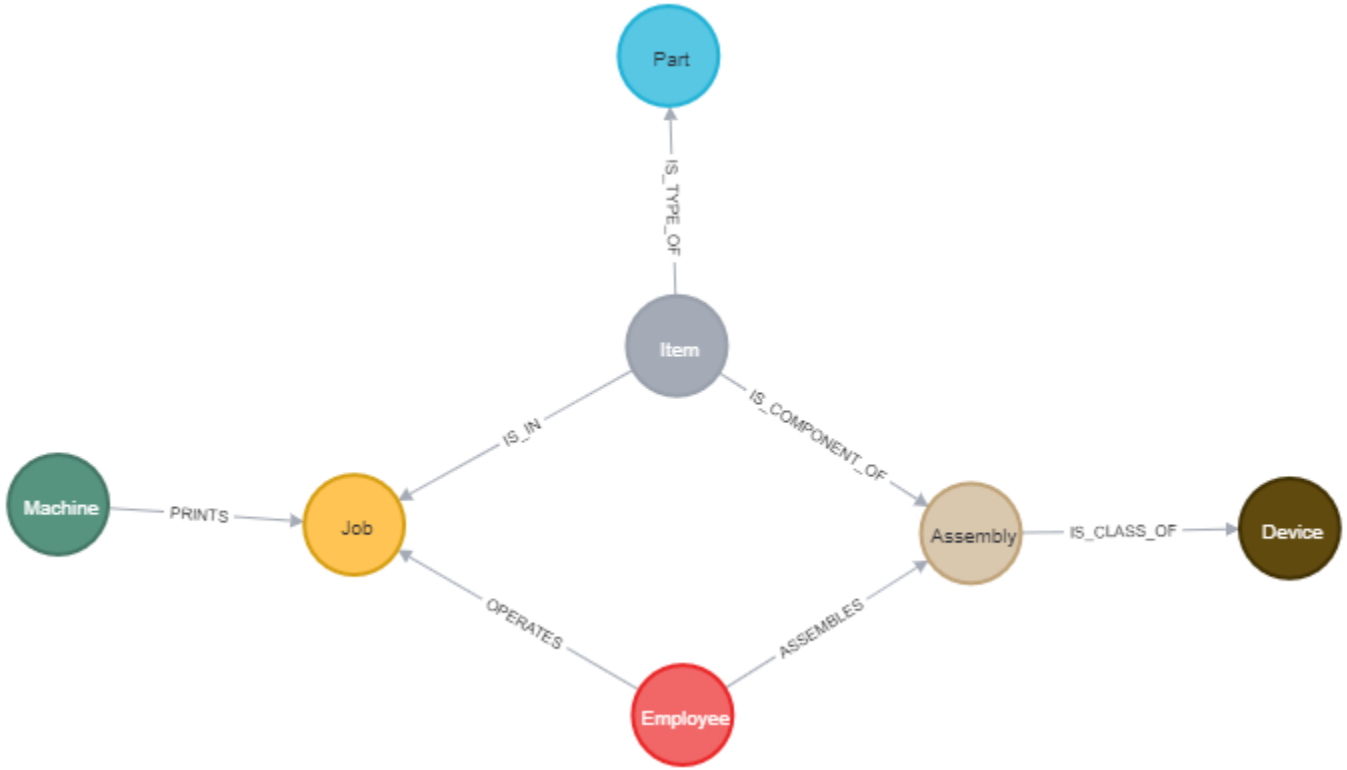

Required queries

1. Query all nodes related to the job_id = 3 Job.
2. Count all parts in job_id = 3.
3. Query all nodes related to the Employee Alice.
4. Count all devices assembled by Employee Chuck.
5. Count all parts printed by the Machine delta.

1. Query all nodes related to job_id = 3

- If we only query the Job nodes associated with jobID='3', we will only get a single node!
- This is NOT what this particularly query was looking for.

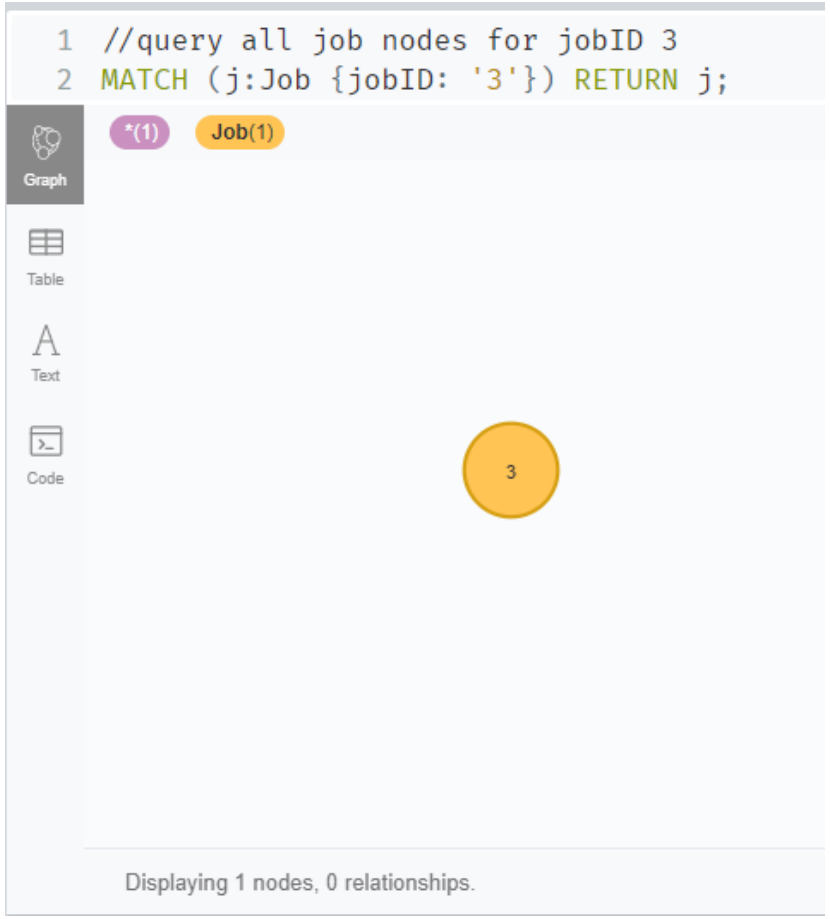

A variable length path makes it easy to query any node label within 1 relationship or hop away from the Job node with jobID = ‘3’

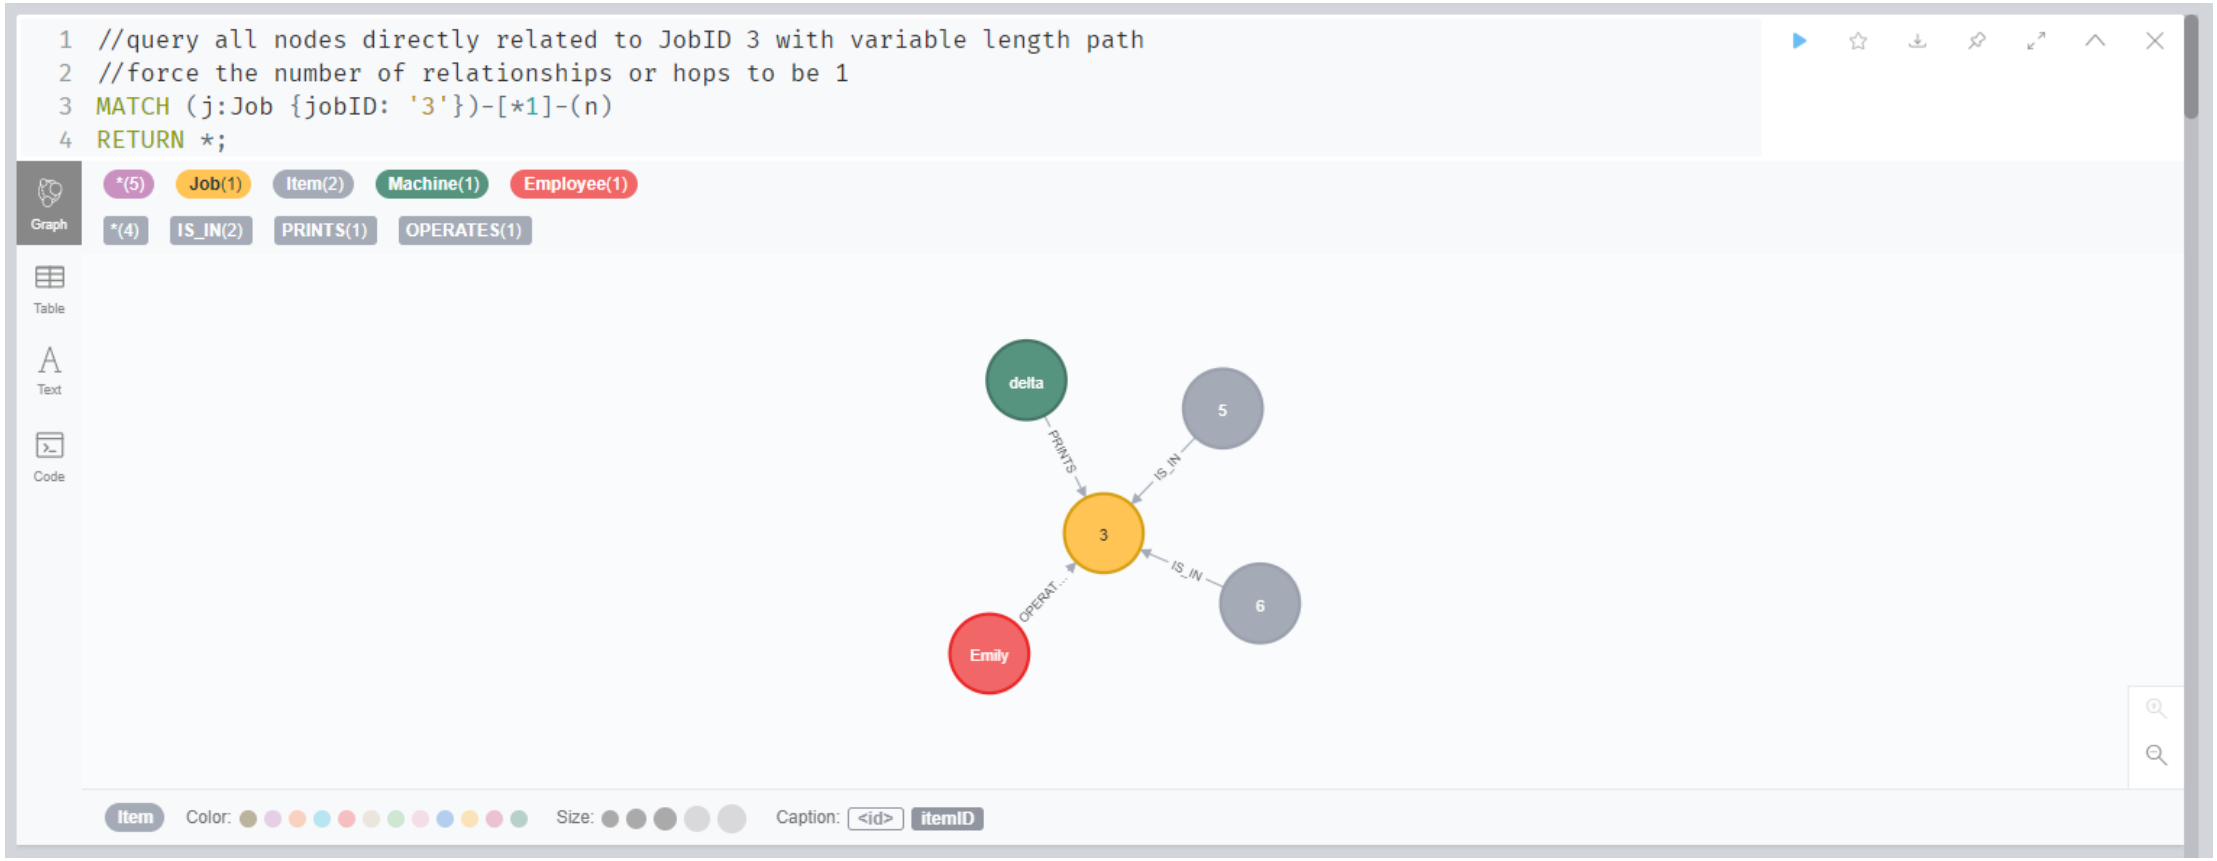

However, we do not need a variable length path for this particular query. We can simply query ANY node Label related in any way to the Job node with jobID='3'

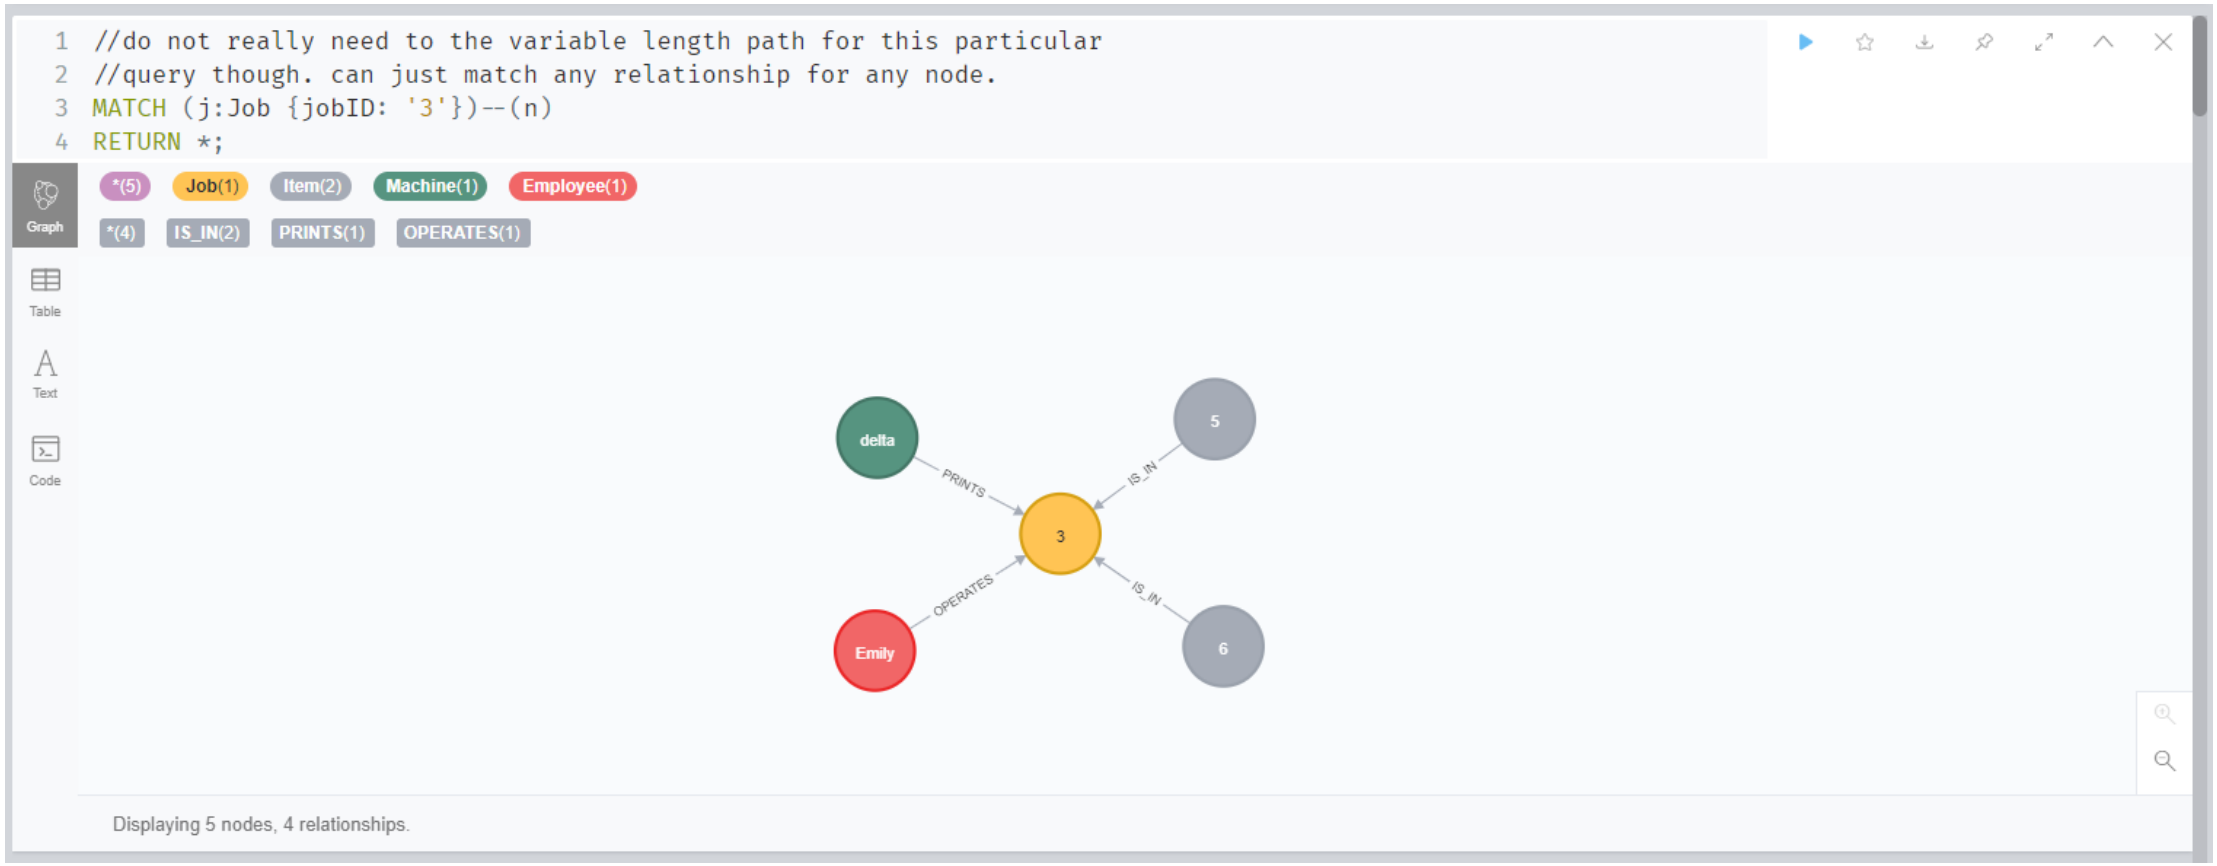

Table view gives the “JSON-like” representation of the query, which is saved as CSV file

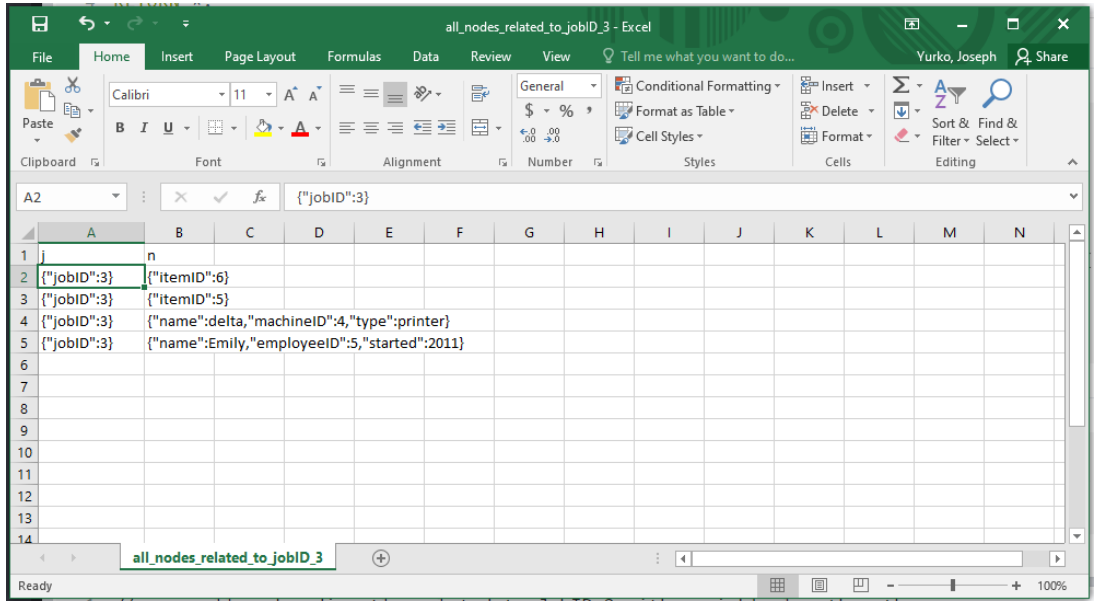

Let’s check our graph query with Pandas

We first need to JOIN all necessary DataFrames together

In [48]:
```
q01 = df_pij.merge( df_mpj, on='job_id', how='left' ).\
merge( df_machines.rename(columns={'name': 'machine_name', 'type': 'machine_type'}), on='machine_id', how='left').\
merge( df_employees.rename(columns={'name': 'employee_name'}), on='employee_id', how='left').\
merge( df_parts.loc[:, ['part_id', 'type']], on='part_id', how='left')
```

In [49]:
```
q01
```

Out[49]:

	item_id	part_id	job_id	machine_id	employee_id	machine_name	machine_type	employee_name	started	type
0	1	4	1	2	2	bravo	printer	Bob	2018	sprocket
1	2	4	1	2	2	bravo	printer	Bob	2018	sprocket
2	3	1	2	4	4	delta	printer	Dave	2019	widget
3	4	3	2	4	4	delta	printer	Dave	2019	gadget
4	5	2	3	4	5	delta	printer	Emily	2011	gizmo
...	...	...	...	...	...	...	...	...	...	...
144	145	5	48	2	5	bravo	printer	Emily	2011	button
145	146	3	48	2	5	bravo	printer	Emily	2011	gadget
146	147	5	49	3	2	charlie	printer	Bob	2018	button
147	148	5	49	3	2	charlie	printer	Bob	2018	button
148	149	3	49	3	2	charlie	printer	Bob	2018	gadget

149 rows × 10 columns

Then, filter or subset to the specific Job of interest

```
In [51]: q01.loc[ q01.job_id == 3, ['job_id', 'employee_name', 'machine_name', 'item_id']]
```

Out[51]:

	job_id	employee_name	machine_name	item_id
4	3	Emily	delta	5
5	3	Emily	delta	6

We get the SAME information as returned by our GRAPH!

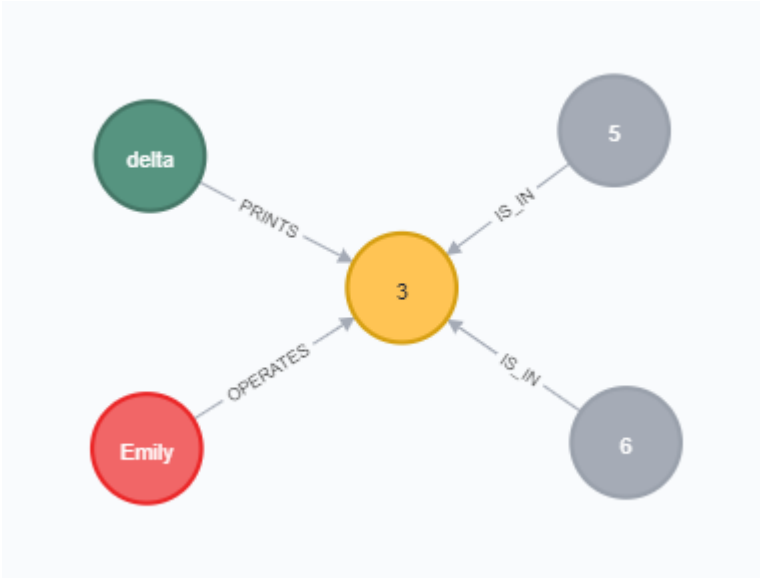

2. Count all parts in job_id = 3.

- The previous query revealed that there are 2 items printed in job_id = 3.

- We can count the number of distinct items using the DISTINCT keyword and count() functions in Cypher.

Query Job and Item nodes related in any way, return the number of unique items per job

```
1 //count all parts for job_id=3
2 //consider the number of unique Items
3 //identify the item node in the query rather than a general node
4 MATCH (j:Job {jobID: '3'})--(item:Item)
5 RETURN j.jobID, count(DISTINCT item.itemID)
```

Table | Text | Code

	j.jobID	count(DISTINCT item.itemID)
1	"3"	2

Started streaming 1 records after 6 ms and completed after 7 ms.

Alternatively, we could count the unique TYPES of parts for job_id = 3.

- Requires including the relationship between the Item node and Part node.

- But allows us to count the number of unique TYPES of parts for a given job.

We get the same result as the previous query which counted the unique number of Items

```
1 //count the unique part types for job_id=3
2 MATCH(j:Job {jobID: '3'})--(item:Item)--(part:Part)
3 RETURN j.jobID, count(DISTINCT part.type);
```

Table | Text | Code

	j.jobID	count(DISTINCT part.type)
1	"3"	2

Started streaming 1 records after 7 ms and completed after 9 ms.

Confirm this is indeed the case with Pandas

As we see below, the two items are two different types of parts!

```
1 //count the unique part types for job_id=3
2 MATCH(j:Job {jobID: '3'})--(item:Item)--(part:Part)
3 RETURN j.jobID, count(DISTINCT part.type);
```

	j.jobID	count(DISTINCT part.type)
1	"3"	2

```
In [52]: q01.loc[ q01.job_id == 3, ['job_id', 'employee_name', 'machine_name', 'item_id', 'type']]
```

Out[52]:

	job_id	employee_name	machine_name	item_id	type
4	3	Emily	delta	5	gizmo
5	3	Emily	delta	6	sprocket

If we do not care about the directionality of the relationship, we do not need to specify it

No arrows are included in the pattern!

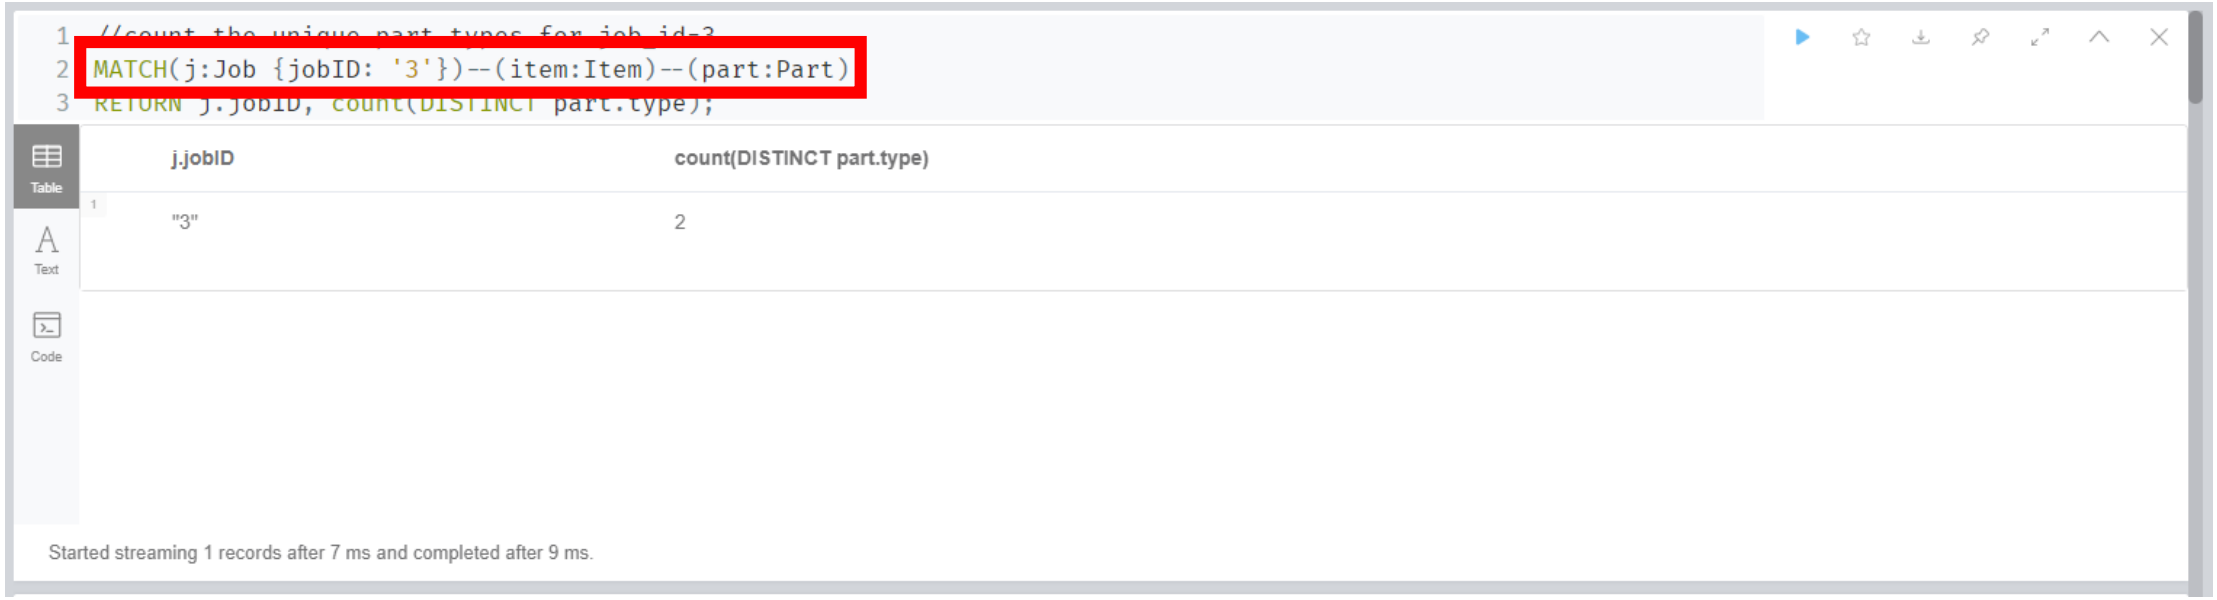

If we want a specific directional relationship, we must be sure that is the direction embedded in the graph

For example, if we give the WRONG directional relationship no results will be returned!

```
1 //count the unique part types for job_id=3
2 //what happens if we include the wrong directions?
3 MATCH(j:Job {jobID: '3'})⟶(item:Item)⟶(part:Part)
4 RETURN j.jobID, count(DISTINCT part.type);
```

Table

Code

(no changes, no records)

Completed after 7 ms.

The query shown above is incorrect because the directionality is Item → Job NOT Job → Item!
As shown by the meta-graph below.

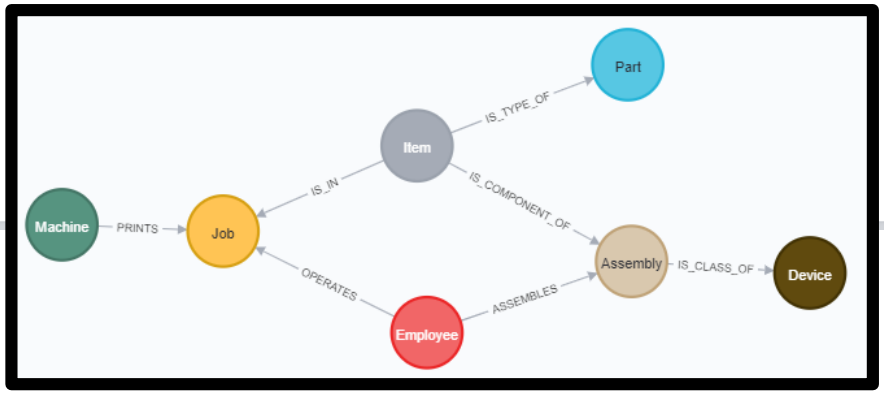

The correct directional query has the arrows pointing away from the Item node

```
1 //count the unique part types for job_id=3
2 //what happens if we include the correct direction?
3 MATCH(j:Job {jobID: '3'})<--(item:Item)-->(part:Part)
4 RETURN j.jobID, count(DISTINCT part.type);
```

Table | Text | Code

	j.jobID	count(DISTINCT part.type)
1	"3"	2

Started streaming 1 records after 6 ms and completed after 7 ms.

3. Query all nodes related to the Employee Alice.

- This query can be structured multiple ways.

One way is to specify any node of any relationship type with the Employee node

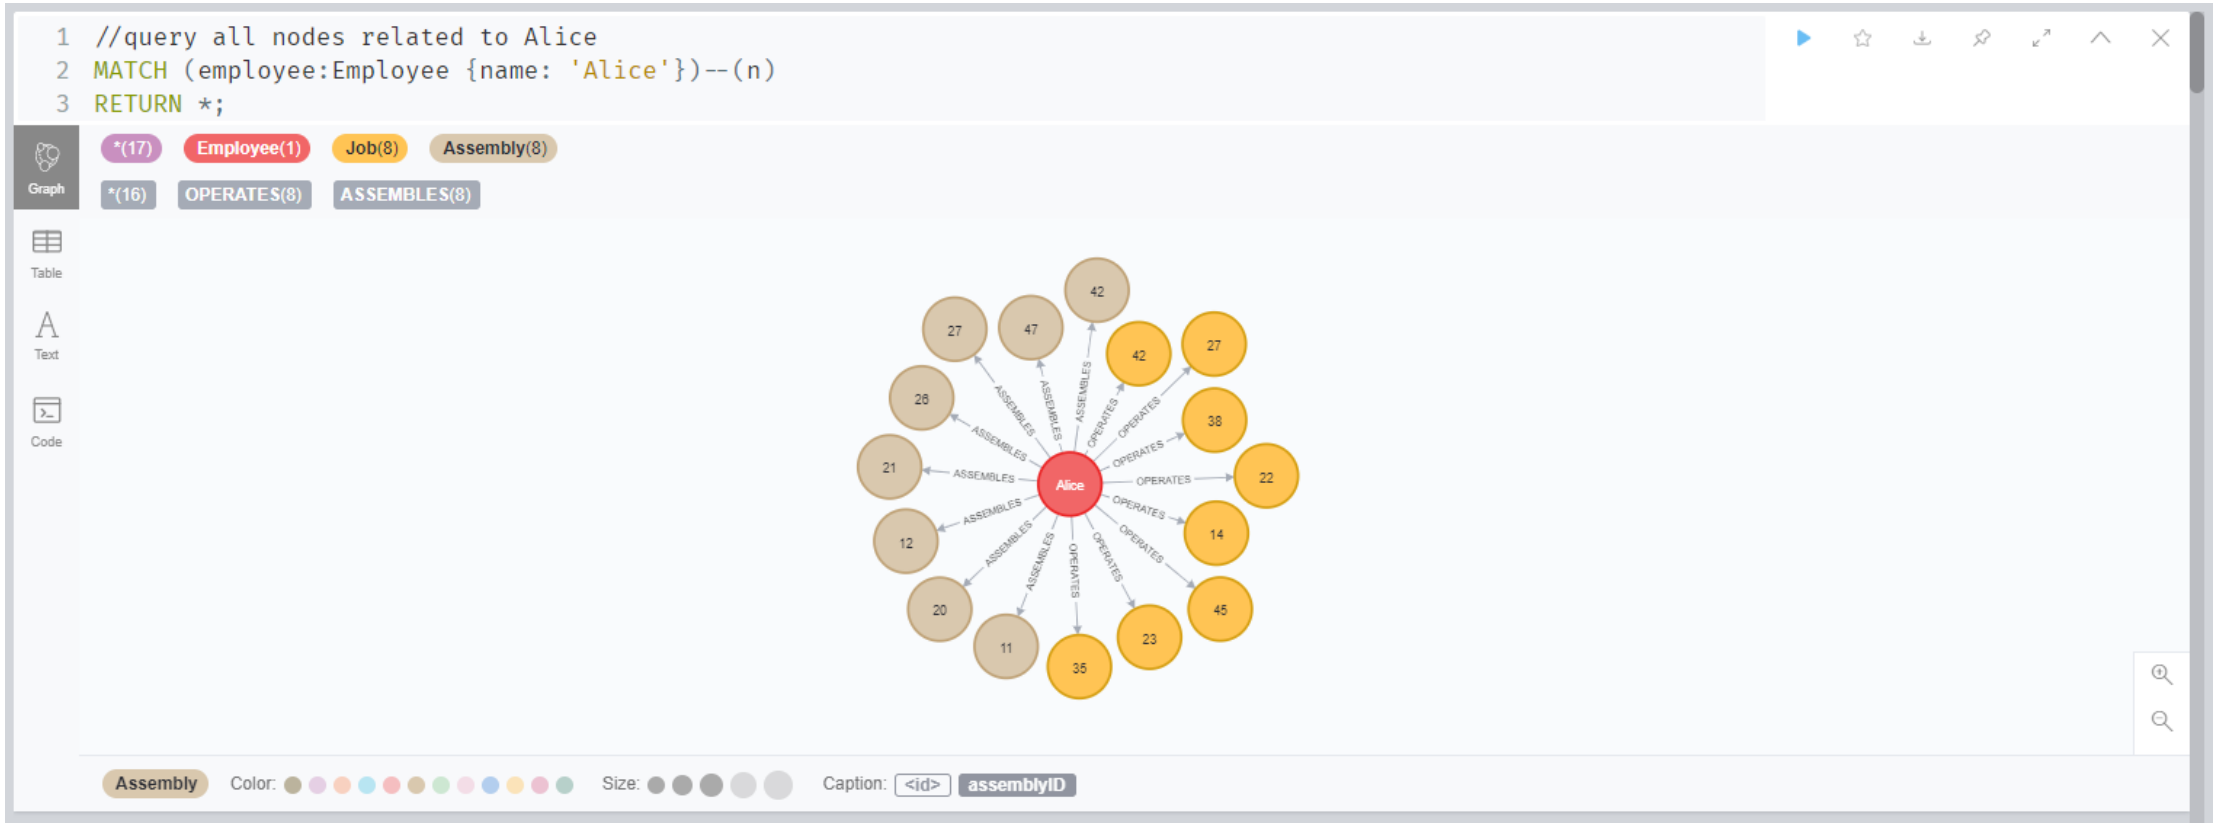

There are 8 Jobs and 8 Assemblies connected with Alice.

To check that this is correct, let's use Pandas to count the number of Jobs which Alice was the machine operator.

```
In [59]: q03_a = df_mpj.merge( df_employees, on='employee_id', how='left' )
```

```
In [60]: q03_a.loc[ q03_a.name == 'Alice', :]
```

Out[60]:

	job_id	machine_id	employee_id	name	started
13	14	4	1	Alice	2017
21	22	3	1	Alice	2017
22	23	3	1	Alice	2017
26	27	3	1	Alice	2017
34	35	1	1	Alice	2017
37	38	2	1	Alice	2017
41	42	2	1	Alice	2017
44	45	3	1	Alice	2017

```
In [62]: q03_a.loc[ q03_a.name == 'Alice', :].shape[0]
```

```
Out[62]: 8
```

Yes! There are 8 jobs!

```
In [63]: q03_b = df_ead.merge( df_employees, on='employee_id', how='left' )
```

```
In [64]: q03_b.loc[ q03_b.name == 'Alice', :]
```

Out[64]:

	assembly_id	device_id	employee_id	name	started
10	11	2	1	Alice	2017
11	12	2	1	Alice	2017
19	20	2	1	Alice	2017
20	21	4	1	Alice	2017
25	26	1	1	Alice	2017
26	27	1	1	Alice	2017
41	42	2	1	Alice	2017
46	47	1	1	Alice	2017

```
In [65]: q03_b.loc[ q03_b.name == 'Alice', :].shape[0]
```

```
Out[65]: 8
```

Yes! There are 8 assemblies!

Alternatively, we can perform the same query by providing the specific node Labels we are interested in.

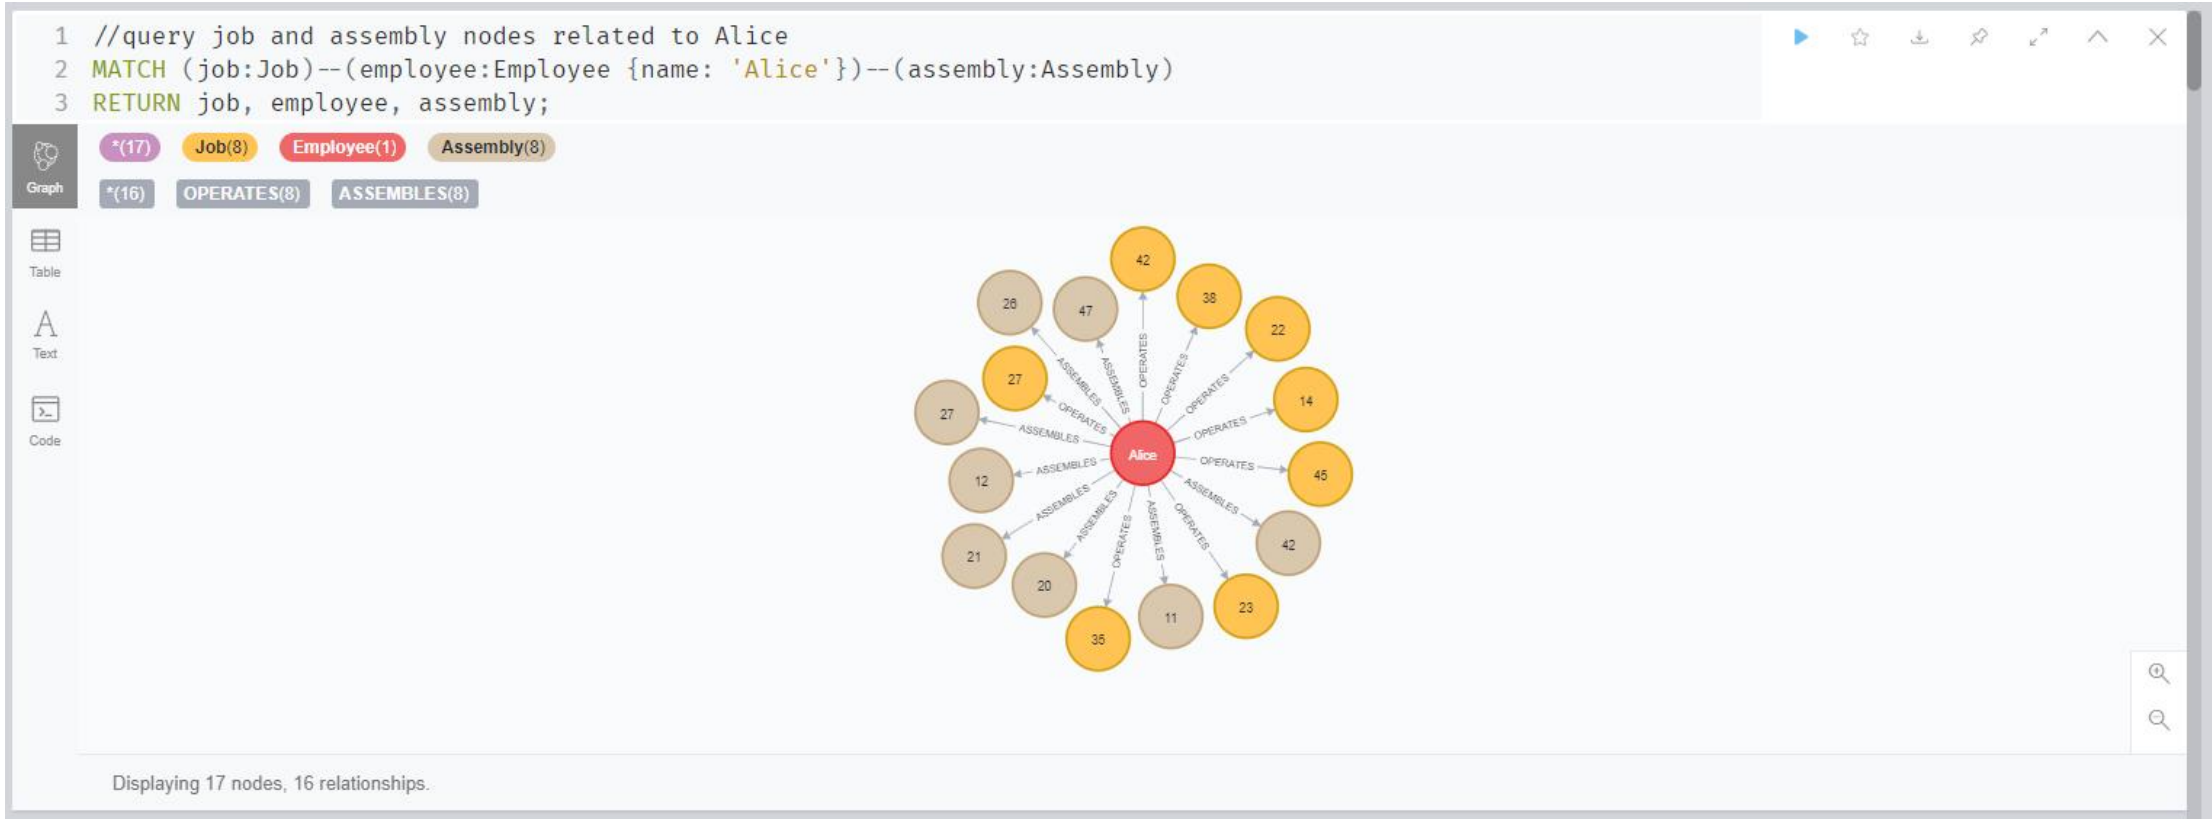

The number of nodes and relationships returned are the same as the previous query.
Though the graphical ordering of the nodes is slightly different.

Alternatively, we can perform the same query by providing the specific node Labels we are interested in.

Easier to structure this query pattern with the Employee node in the middle.

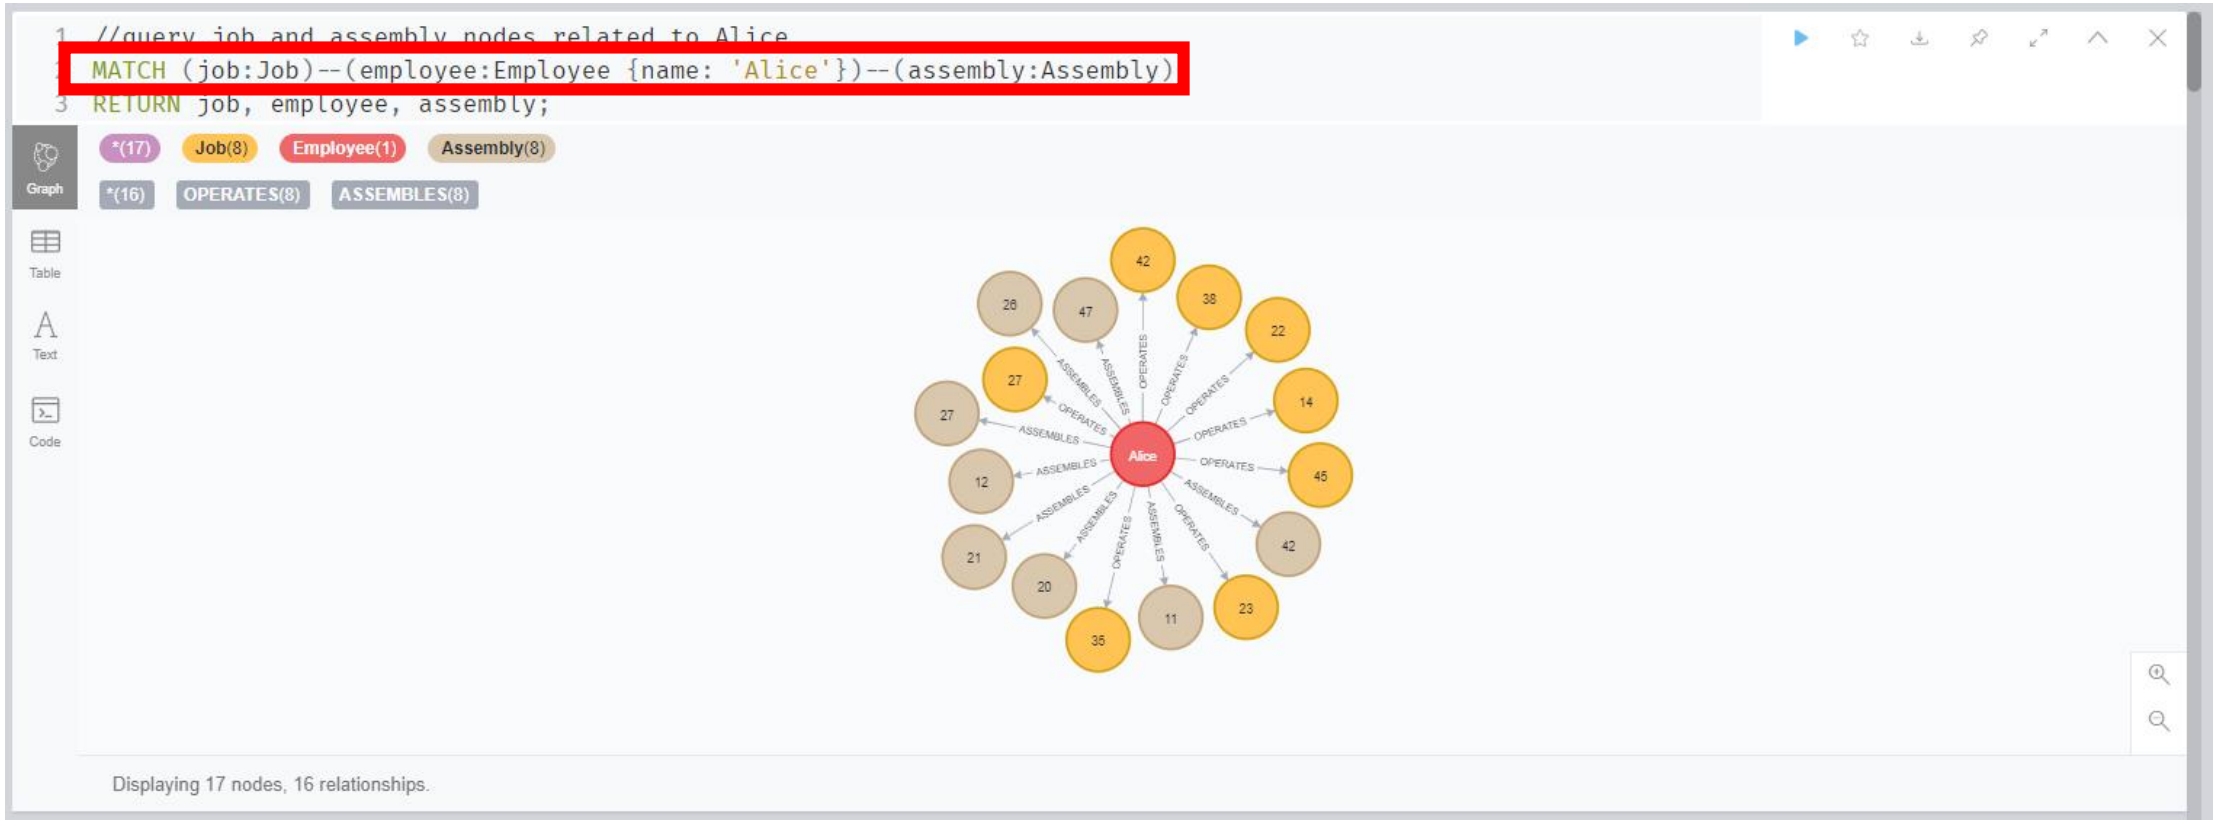

The number of nodes and relationships returned are the same as the previous query.
Though the graphical ordering of the nodes is slightly different.

4. Count all devices assembled by Employee Chuck

- We can count the number of assemblies by querying all Assemblies related to the Employee node with name='Chuck'.
- If we want to also count the unique number of device classes, we will need to relate the Assembly node to the Device node.
- The question itself was open ended, you could have answered either way!

Querying all Assembly nodes related to Chuck returns 13 Assembly nodes

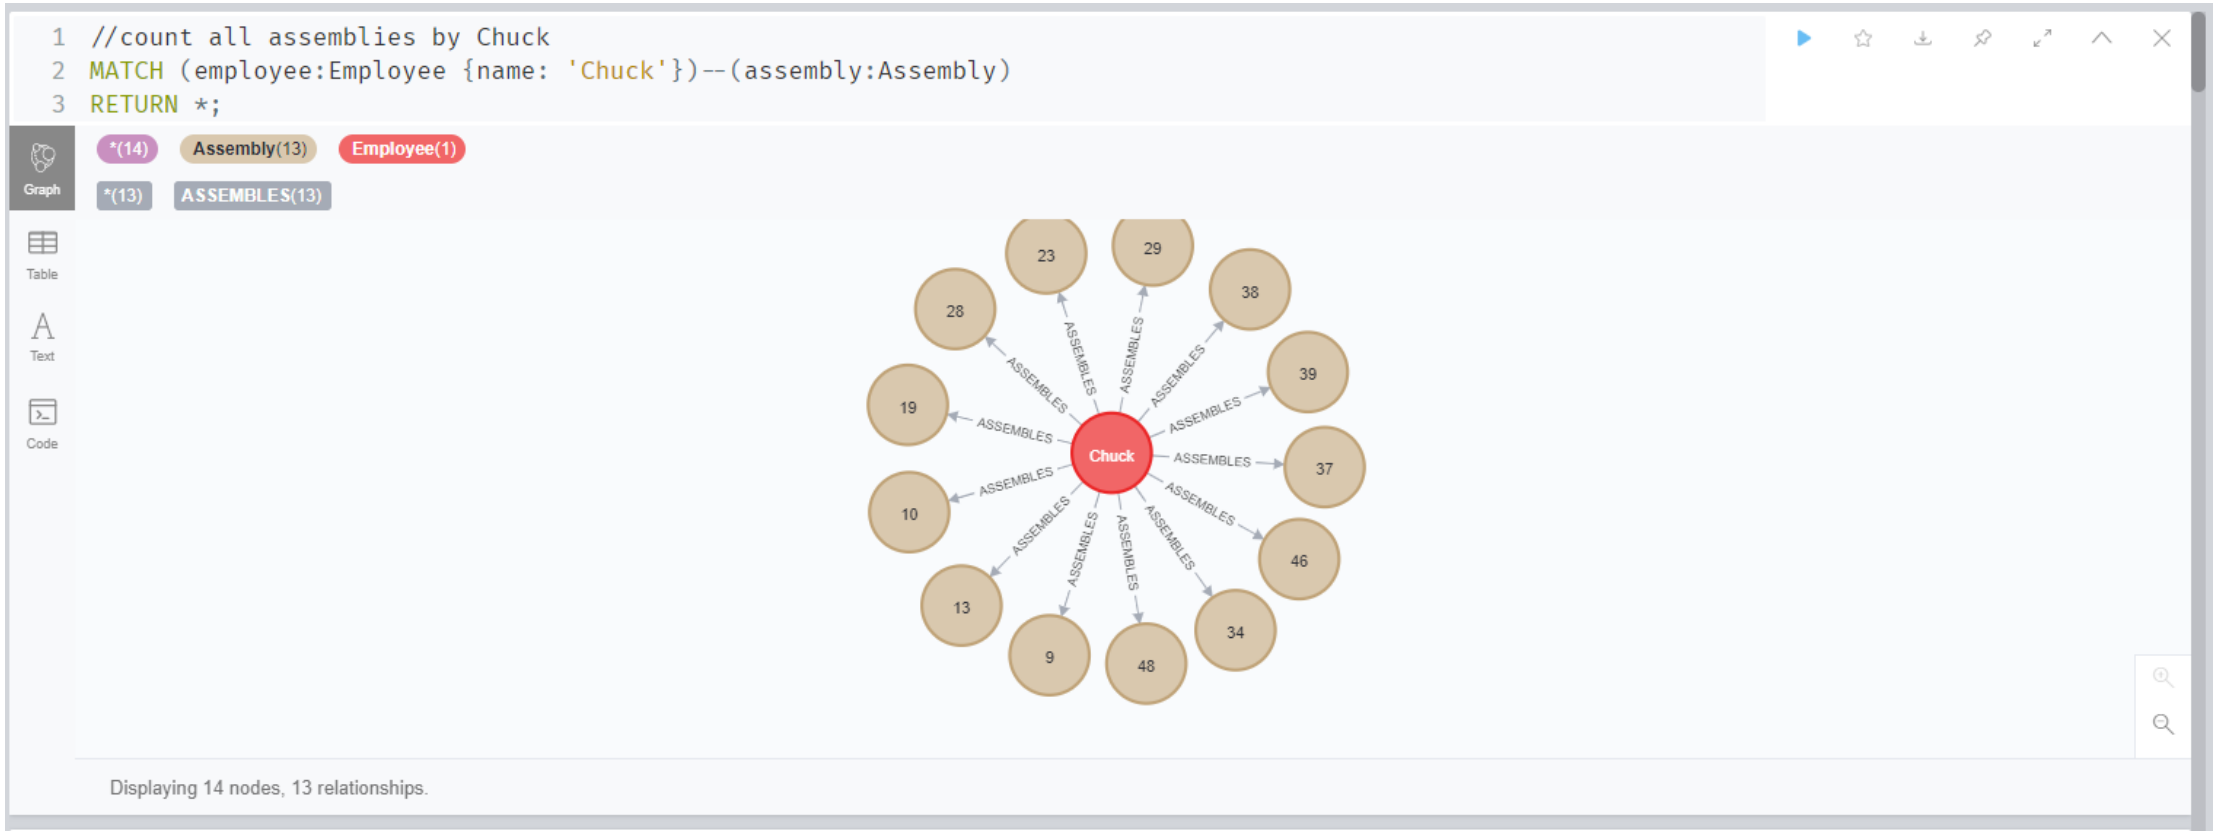

Perform the grouping and aggregation operation with the count() function

```
1 //count all assemblies by Chuck
2 MATCH (employee:Employee {name: 'Chuck'})--(assembly:Assembly)
3 RETURN employee.name, count(assembly.assemblyID);
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)
1	"Chuck"	13

Started streaming 1 records after 6 ms and completed after 8 ms.

We get the same count we could visually confirm from the previous query.

Confirm by counting the rows in Pandas

```
1 //count all assemblies by Chuck
2 MATCH (employee:Employee {name: 'Chuck'})--(assembly:Assembly)
3 RETURN employee.name, count(assembly.assemblyID);
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)
1	"Chuck"	13

```
In [67]: q03_b.loc[ q03_b.name == 'Chuck', :].shape[0]
Out[67]: 13
```

Started streaming 1 records after 6 ms and completed after 8 ms.

Or by grouping and aggregating

```
1 //count all assemblies by Chuck
2 MATCH (employee:Employee {name: 'Chuck'})--(assembly:Assembly)
3 RETURN employee.name, count(assembly.assemblyID);
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)
1	"Chuck"	13

Started streaming 1 records after 6 ms and completed afte

```
In [70]: q03_b.groupby(['name']).\
         aggregate(num_assemblies = ('assembly_id', 'nunique')).\
         reset_index()
```

Out[70]:

	name	num_assemblies
0	Alice	8
1	Bob	10
2	Chuck	13
3	Dave	8
4	Emily	11

To count the number of unique Device classes assembled by Chuck:

- We need to include the relationship between the Assembly node and the Device node.

- Apply the count() function to the DISTINCT device.class property.

Notice that the pattern includes 3 nodes and 2 relationships now

```
1 //count all assemblies and unique devices by Chuck
2 MATCH (employee:Employee {name: 'Chuck'})--(assembly:Assembly)--(device:Device)
3 RETURN employee.name, count(assembly.assemblyID), count(DISTINCT device.class);
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)	count(DISTINCT device.class)
1	"Chuck"	13	4

Started streaming 1 records after 12 ms and completed after 14 ms.

Confirm by grouping and aggregating in Pandas

```
1 //count all assemblies and unique devices by Chuck
2 MATCH (employee:Employee {name: 'Chuck'})--(assembly:Assembly)--(device:Device)
3 RETURN employee.name, count(assembly.assemblyID), count(DISTINCT device.class);
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)	count(DISTINCT device.class)
1	"Chuck"	13	4

Started streaming 1 records after 12 ms and completed after 14 ms.

```
In [71]: q03_b.groupby(['name']).\
         aggregate(num_assemblies = ('assembly_id', 'nunique'),
                   num_device_classes = ('device_id', 'nunique')).\
         reset_index()
```

Out[71]:

	name	num_assemblies	num_device_classes
0	Alice	8	3
1	Bob	10	4
2	Chuck	13	4
3	Dave	8	4
4	Emily	11	4

Although not required, we could count the number of Assemblies and unique Device classes for every employee in Cypher

```
1 //count all assemblies and unique devices by each employee
2 //not required to do in the question
3 MATCH (employee:Employee)--(assembly:Assembly)--(device:Device)
4 RETURN employee.name, count(assembly.assemblyID), count(DISTINCT device.class)
5 ORDER BY employee.name;
```

Table | Text | Code

	employee.name	count(assembly.assemblyID)	count(DISTINCT device.class)
1	"Alice"	8	3
2	"Bob"	10	4
3	"Chuck"	13	4
4	"Dave"	8	4
5	"Emily"	11	4

Started streaming 5 records after 5 ms and completed after 32 ms.

The Cypher query gives the same results as the Pandas merge(), groupby(), and aggregate() series of statements. Analogous to joining, grouping, and summarizing in MySQL.

```
//count all assemblies and unique devices by each employee
//not required to do in the question
MATCH (employee:Employee)--(assembly:Assembly)--(device:Device)
RETURN employee.name, count(assembly.assemblyID), count(DISTINCT device.class)
ORDER BY employee.name;
```

	employee.name	count(assembly.assemblyID)	count(DISTINCT device.class)
1	"Alice"	8	3
2	"Bob"	10	4
3	"Chuck"	13	4
4	"Dave"	8	4
5	"Emily"	11	4

Started streaming 5 records after 5 ms and completed after 32 ms.

```
In [72]: df_ead.merge( df_employees, on='employee_id', how='left' ).\
groupby(['name']).\
aggregate(num_assemblies = ('assembly_id', 'nunique'),
          num_device_classes = ('device_id', 'nunique')).\
reset_index()
```

Out[72]:

	name	num_assemblies	num_device_classes
0	Alice	8	3
1	Bob	10	4
2	Chuck	13	4
3	Dave	8	4
4	Emily	11	4

5. Count all parts printed by the Machine delta.

- Question is open ended, so could count the unique number of Items printed OR the unique number of Part types.

Although not required, the network for all Jobs printed by Machine delta and all items in those Jobs

The graph tells us there are 12 Jobs and 24 Items.

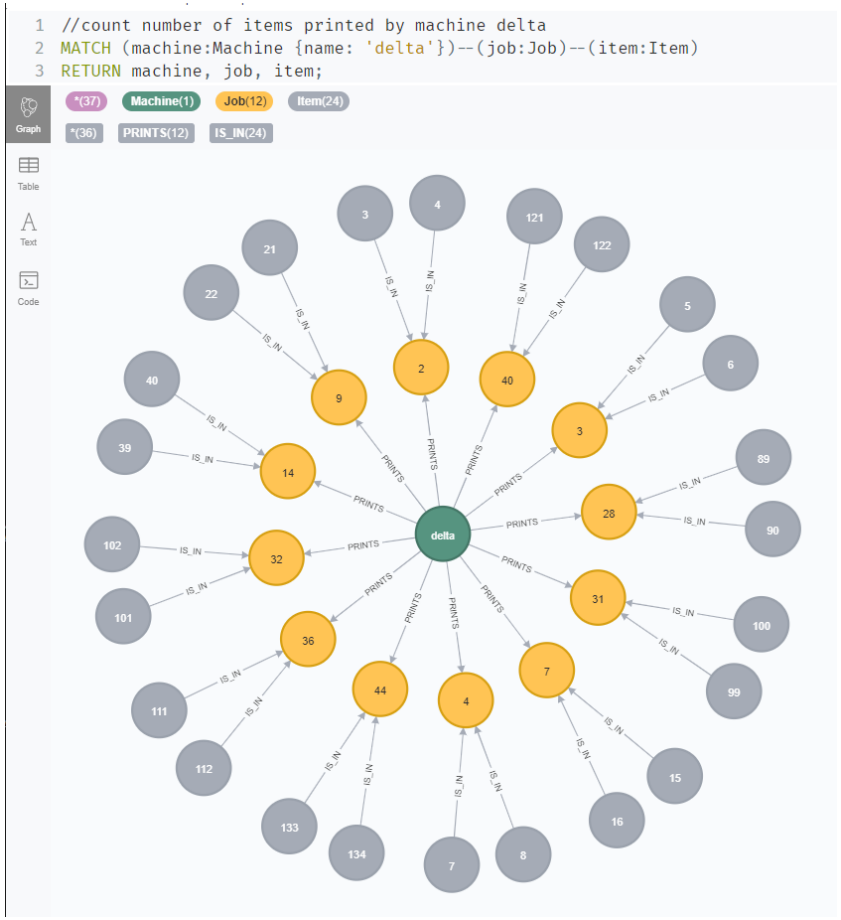

Query the number of Items by grouping and summarizing via the count() function

```
1 //count number of items printed by machine delta
2 MATCH (machine:Machine {name: 'delta'})--(job:Job)--(item:Item)
3 RETURN machine.name, count(DISTINCT job.jobID), count(DISTINCT item.itemID);
```

Table | Text | Code

	machine.name	count(DISTINCT job.jobID)	count(DISTINCT item.itemID)
1	"delta"	12	24

Started streaming 1 records after 7 ms and completed after 9 ms.

The number of unique Jobs is included, but was not required for the question.

If we also want to query the number of unique Part types printed by Machine delta, we need to include the Item to Part relationship and apply the count() function appropriately

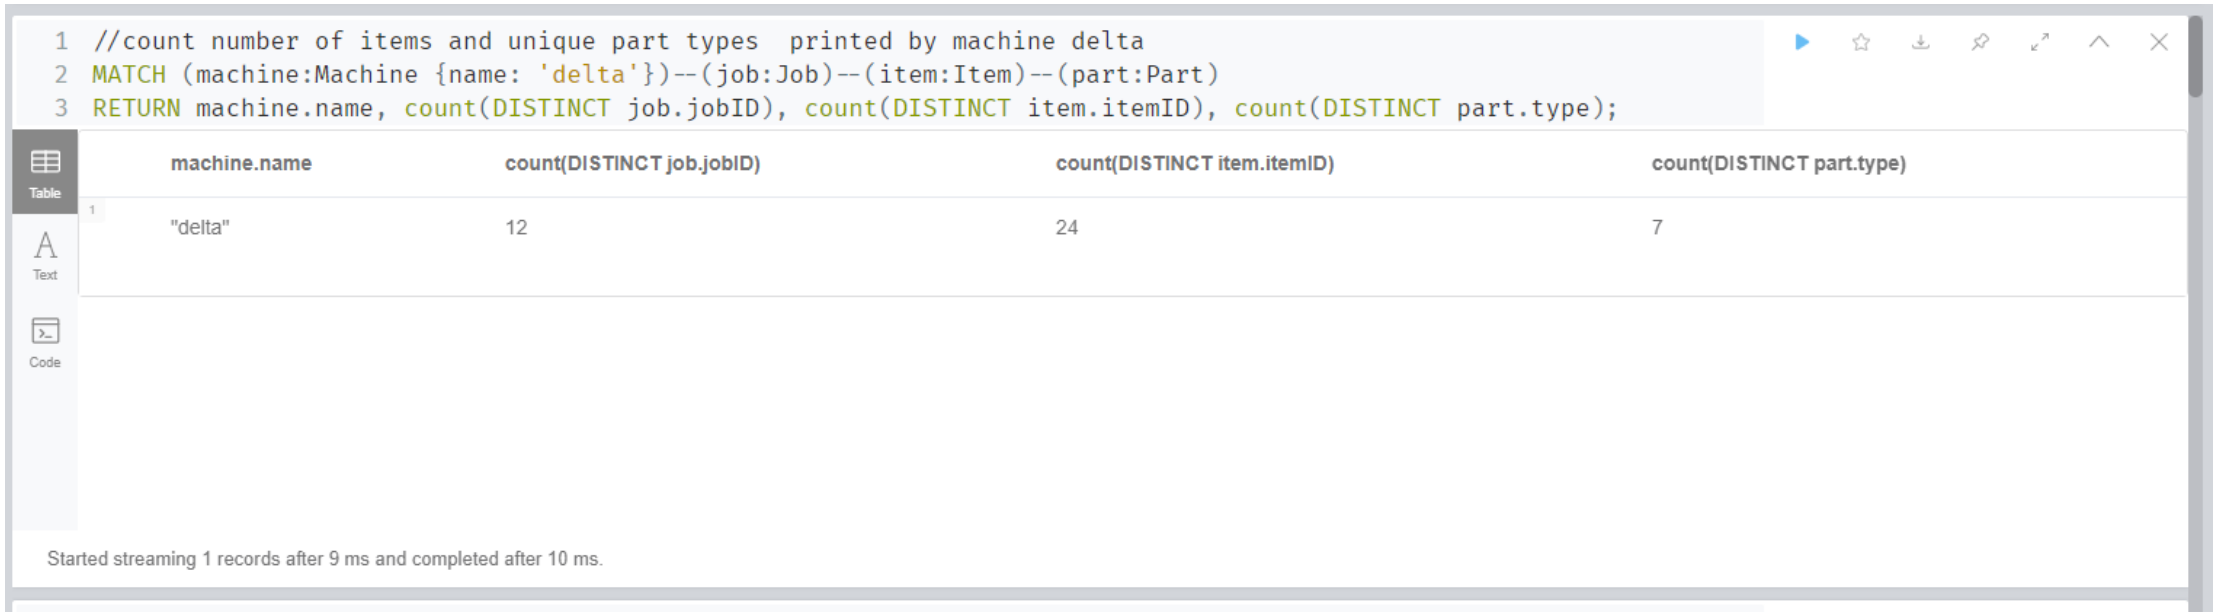

Although not required, it's straight forward to repeat the query for all machines not just Machine delta

```
1 //count number of items and unique part types printed for all machines
2 MATCH (machine:Machine)--(job:Job)--(item:Item)--(part:Part)
3 RETURN machine.name, count(DISTINCT job.jobID), count(DISTINCT item.itemID), count(DISTINCT part.type)
4 ORDER BY machine.name;
```

Table | Text | Code

	machine.name	count(DISTINCT job.jobID)	count(DISTINCT item.itemID)	count(DISTINCT part.type)
1	"alpha"	12	48	7
2	"bravo"	11	22	5
3	"charlie"	14	55	7
4	"delta"	12	24	7

Started streaming 4 records after 8 ms and completed after 15 ms.

Check the results with Pandas

```
1 //count number of items and unique part types printed for all machines
2 MATCH (machine:Machine)--(job:Job)--(item:Item)--(part:Part)
3 RETURN machine.name, count(DISTINCT job.jobID), count(DISTINCT item.itemID), count(DISTINCT part.type)
4 ORDER BY machine.name;
```

Table | Text | Code

	machine.name	count(DISTINCT job.jobID)	count(DISTINCT item.itemID)	count(DISTINCT part.type)
1	"alpha"	12	48	7
2	"bravo"	11	22	5
3	"charlie"	14	55	7
4	"delta"	12	24	7

Started streaming 4 records after 8 ms and co

```
In [80]: df_pij.merge( df_mpj, on='job_id', how='left').\
         merge( df_machines.loc[:, ['machine_id', 'name']], on='machine_id', how='left').\
         merge( df_parts, on='part_id', how='left').\
         groupby(['name']).\
         aggregate(num_jobs = ('job_id', 'nunique'),
                   num_items = ('item_id', 'nunique'),
                   num_part_types = ('type', 'nunique')).\
         reset_index()
```

Out[80]:

	name	num_jobs	num_items	num_part_types
0	alpha	12	48	7
1	bravo	11	22	5
2	charlie	14	55	7
3	delta	12	24	7